

```
In [83]: # numerical analysis
import numpy as np
import pandas as pd
from pandas.plotting import autocorrelation_plot

# import statistics package
import statsmodels
import statsmodels.api as sm
import statsmodels.tsa.api as smt
from statsmodels.tsa.stattools import grangercausalitytests
from statsmodels.stats.diagnostic import het_arch
from statsmodels.compat import lzip
# from arch import arch_model

# visualisation
from matplotlib import pyplot as plt
import matplotlib.mlab as mlab
import seaborn as sns
import plotly.express as px

# machine learning
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.tree import DecisionTreeClassifier, plot_tree, DecisionTreeRegressor
from sklearn.svm import SVC, SVR
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier, GradientBoostingRegressor
from sklearn.neural_network import MLPClassifier, MLPRegressor
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score
```

Part 1 - Descriptive Statistics and Visuals

Q1: Describe both banks individually using any appropriate descriptive statistics and visuals, including ARMA processes and unit roots

```
In [84]: prices = pd.read_csv('IrishBanks.csv', index_col=0)
prices.index=pd.to_datetime(prices.index, infer_datetime_format=True)
print(prices.head(3))
print(prices.tail(3))

# Data starts 02 January 2018
# Ends 05 February 2021
# Both stock prices have reduced significantly
```

	AIB	BoI
Date		
2018-01-02	5.460	7.195
2018-01-03	5.435	7.370
2018-01-04	5.450	7.545
	AIB	BoI
Date		
2021-02-03	1.542	3.326
2021-02-04	1.554	3.290
2021-02-05	1.584	3.282

```
In [85]: prices.describe()

# At its peak, AIB reached €5.80, BoI reached €8.15
# The lowest they fell to was €0.76 and €1.33
```


Out [85]:

	AIB	BoI
count	789.000000	789.000000
mean	3.132753	4.642466
std	1.497407	2.021505
min	0.768000	1.330000
25%	1.554000	3.020000
50%	3.498000	4.700000
75%	4.424000	6.480000
max	5.800000	8.150000


```
In [86]: fig1= px.line(prices, x= prices.index, y= "AIB", title = 'AIB stock
fig2= px.line(prices, x= prices.index, y= "BoI", title = 'BoI stock
fig1.show()
fig2.show()

# AIB and BoI prices move in a very similar downward trend
# These series appears to be non-stationary
```

AIB stock price over time




BoI stock price over time



```
In [87]: plt.plot(prices.index, prices['AIB'], color='red', label = 'AIB')
plt.plot(prices.index, prices['BoI'], color='blue', label = 'BoI')
plt.xlabel('Time')
plt.ylabel('Stock Prices')
plt.title('Stock prices over time')
plt.legend()
plt.show()
```

The similarity can be seen clearly below



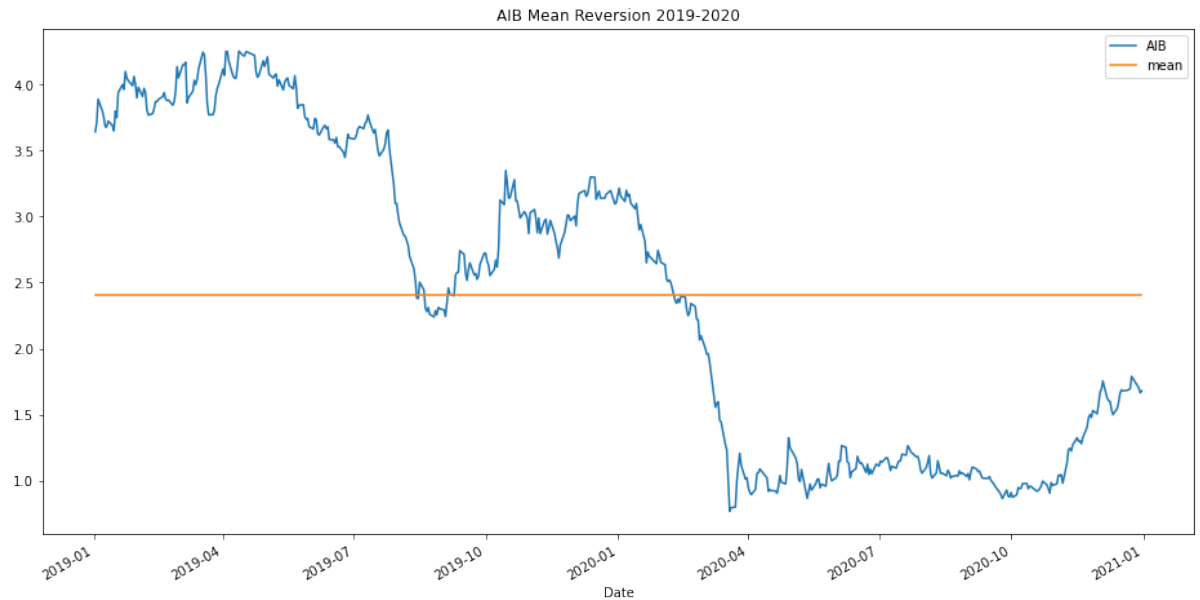
Are the stocks mean reverting?

```
In [88]: AIB19_20 = prices.loc['2019':'2020']['AIB']
```

```
In [89]: AIB19_20['mean'] = AIB19_20['AIB'].mean()
```

```
In [90]: AIB19_20.plot(figsize=(16,8), title='AIB Mean Reversion 2019-2020')  
# It does not look like AIB's stock price reverts to the mean
```

```
Out[90]: <AxesSubplot:title={'center':'AIB Mean Reversion 2019-2020'}, xlabel='Date'>
```




```
In [91]: # Now test BoI  
BoI19_20 = prices.loc['2019':'2020'][['BoI']]
```

```
In [92]: BoI19_20['mean'] = BoI19_20['BoI'].mean()
```

```
In [93]: BoI19_20.plot(figsize=(16,8), title='BoI Mean Reversion 2019-2020')
# BoI's stock price looks slightly more mean reverting
```

```
Out[93]: <AxesSubplot:title={'center':'BoI Mean Reversion 2019-2020'}, xlabel='Date'>
```



```
In [94]: # Import the ISEQ as a comparison to the overall Irish stock exchange

ISEQ = pd.read_csv('ISEQ.csv', index_col=0)
ISEQ.index=pd.to_datetime(ISEQ.index, infer_datetime_format=True)
print(ISEQ.head(3))
print(ISEQ.tail(3))
```


	ISEQ
Date	
2018-01-02	7054.549805
2018-01-03	7071.229980
2018-01-04	7126.569824
Date	
2021-02-03	7337.609863
2021-02-04	7358.020020
2021-02-05	7361.959961

```
In [95]: # How do BoI returns compare to the ISEQ?


fig1= px.line(prices, x= prices.index, y= "BoI", title = 'BoI price')
fig2= px.line(ISEQ, x= ISEQ.index, y= "ISEQ", title = 'ISEQ price')
fig1.show()
fig2.show()

# BoI and the ISEQ have somewhat of a similar pattern
```

BoI price over time



ISEQ price over time




In [96]: *# Show the distribution of prices*

```
plt.hist(prices['AIB'], bins=50, alpha = 0.4)
plt.hist(prices['BoI'], bins=50, alpha = 0.4)
```

Very similar patterns

Out[96]: (array([11., 13., 35., 61., 33., 11., 3., 4., 3., 8., 5., 8.,
7.,
14., 20., 23., 16., 9., 9., 8., 10., 13., 19., 31., 28.,
24.,
22., 15., 17., 19., 19., 20., 15., 12., 2., 8., 12., 6.,
8.,
20., 8., 10., 27., 26., 38., 21., 10., 14., 9., 5.]),
array([1.33 , 1.4664, 1.6028, 1.7392, 1.8756, 2.012 , 2.1484, 2.
2848,
2.4212, 2.5576, 2.694 , 2.8304, 2.9668, 3.1032, 3.2396, 3.
376 ,
3.5124, 3.6488, 3.7852, 3.9216, 4.058 , 4.1944, 4.3308, 4.
4672,
4.6036, 4.74 , 4.8764, 5.0128, 5.1492, 5.2856, 5.422 , 5.
5584,
5.6948, 5.8312, 5.9676, 6.104 , 6.2404, 6.3768, 6.5132, 6.
6496,
6.786 , 6.9224, 7.0588, 7.1952, 7.3316, 7.468 , 7.6044, 7.
7408,
7.8772, 8.0136, 8.15]),
<BarContainer object of 50 artists>)



In [97]: `prices.corr()`

98% correlation between AIB and BoI stock prices

Out [97]:

	AIB	BoI
AIB	1.00000	0.98398
BoI	0.98398	1.00000

Checking Stationarity of Stock Prices

In [98]: `test1 = prices.loc['2020-02-05']`
`train1 = prices.loc[:'2020-02-04']`


In [99]: *# Testing stationarity of AIB Stock Price*
from statsmodels.tsa.stattools **import** adfuller

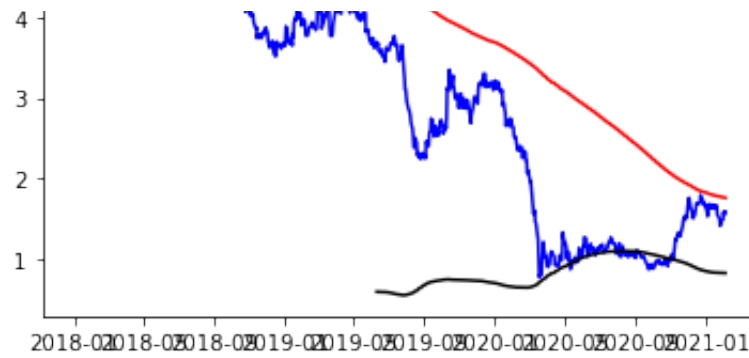
```
def test_stationarity(timeseries):
    #Determining rolling statistics
    rolmean = prices['AIB'].rolling(window=365).mean()
    rolstd = prices['AIB'].rolling(window=365).std()
    #Plot rolling statistics:
    plt.plot(prices['AIB'], color='blue', label='Original')
    plt.plot(rolmean, color='red', label='Rolling Mean')
    plt.plot(rolstd, color='black', label='Rolling Std')
    plt.legend(loc='best')
    plt.title('Rolling Mean and Standard Deviation')
    plt.show(block=False)

    print("Results of dickey fuller test")
    adft = adfuller(timeseries, autolag='AIC')
    # output for dft will give us without defining what the values are
    #hence we manually write what values does it explains using a for
    output = pd.Series(adft[0:4], index=['Test Statistics', 'p-value', 'N', 'L'])
    for key, values in adft[4].items():
        output['critical value (%)'%key] = values
    print(output)

test_stationarity(train1['AIB'])
```

p-value > 0.05, so AIB stock prices are non-stationary
This was to be expected
Rolling Mean and Rolling Standard Deviation show the non-stationary





Results of dickey fuller test

Test Statistics	-0.974615
p-value	0.762410
No. of lags used	0.000000
Number of observations used	531.000000
critical value (1%)	-3.442725
critical value (5%)	-2.866998
critical value (10%)	-2.569677
dtype: float64	


```
In [100]: # Examine the seasonality of AIB prices
from statsmodels.tsa.seasonal import seasonal_decompose
decomposition = seasonal_decompose(prices.AIB, freq = 260) #Avg around 2 years
trend = decomposition.trend
seasonal = decomposition.seasonal
residual = decomposition.resid

plt.figure(figsize=(12,6))
plt.subplot(411)
plt.plot(prices['AIB'], label='Original')
plt.legend(loc='best')
plt.subplot(412)
plt.plot(trend, label='Trend')
plt.legend(loc='best')
plt.subplot(413)
plt.plot(seasonal, label='Seasonality')
plt.legend(loc='best')
plt.subplot(414)
plt.plot(residual, label='Residuals')
plt.legend(loc='best')
plt.tight_layout()
plt.show()

# A clear downward trend in stock prices
```

/var/folders/mh/j7bgd0wx1mb1jjnzcp0c8dfw0000gn/T/ipykernel_83727/842978463.py:3: FutureWarning:

the 'freq' keyword is deprecated, use 'period' instead




```
In [101]: # Examine the seasonality of BoI
from statsmodels.tsa.seasonal import seasonal_decompose
decomposition = seasonal_decompose(prices.BoI, freq = 260) #Avg around 200 days
trend = decomposition.trend
seasonal = decomposition.seasonal
residual = decomposition.resid

plt.figure(figsize=(12,6))
plt.subplot(411)
plt.plot(prices['BoI'], label='Original')
plt.legend(loc='best')
plt.subplot(412)
plt.plot(trend, label='Trend')
plt.legend(loc='best')
plt.subplot(413)
plt.plot(seasonal, label='Seasonality')
plt.legend(loc='best')
plt.subplot(414)
plt.plot(residual, label='Residuals')
plt.legend(loc='best')
plt.tight_layout()
plt.show()
```

/var/folders/mh/j7bgd0wx1mb1jjnzcp0c8dfw0000gn/T/ipykernel_83727/162248132.py:3: FutureWarning:

the 'freq' keyword is deprecated, use 'period' instead



In [102]: *# Transform stock prices to returns*

```
prices['dAIB'] = prices['AIB'].transform(lambda x : (x - x.shift(1)))
prices['dBoI'] = prices['BoI'].transform(lambda x : (x - x.shift(1)))
prices = prices.dropna()
prices.head(3)

# Transforming to returns may turn the series stationary
```


Out[102]:

	AIB	BoI	dAIB	dBoI
Date				
2018-01-03	5.435	7.370	-0.457875	2.432245
2018-01-04	5.450	7.545	0.275989	2.374491
2018-01-05	5.450	7.735	0.000000	2.518224

In [103]: `fig1= px.line(prices, x= prices.index, y= "dAIB", title = 'AIB stock returns over time')`
`fig2= px.line(prices, x= prices.index, y= "dBoI", title = 'BoI stock returns over time')`
`fig1.show()`
`fig2.show()`


Both sets of returns look to be stationary, but have some nonconstant variance
There is no evident upward or downward trend over time

AIB stock returns over time




BoI stock returns over time

BOI STOCK RETURNS OVER TIME



```
In [104]: plt.scatter(prices['dAIB'], prices['dBoI'])  
plt.title('Scatterplot of AIB and BoI Returns')  
plt.show()
```

```
# Positive Correlation is evident  
# When one goes up, the other likely follows
```



```
In [105]: correlation = prices['dAIB'].corr(prices['dBoI'])  
print("Correlation of the stock returns is: ", correlation)
```

Correlation of the stock returns is: 0.7364661575240155

```
In [106]: prices.describe()
```

```
Out[106]:
```

	AIB	Bol	dAIB	dBol
count	788.000000	788.000000	788.000000	788.000000
mean	3.129800	4.639227	-0.093628	-0.045353
std	1.496057	2.020739	3.558539	3.292536
min	0.768000	1.330000	-23.960396	-15.714286
25%	1.552250	3.016000	-1.762296	-1.732064
50%	3.492000	4.698000	-0.181295	-0.144521
75%	4.413500	6.476250	1.250806	1.508710
max	5.800000	8.150000	24.625000	15.458015

Unit Root Tests

```
In [107]: acf_AIB = smt.acf(prices['dAIB'], nlags=15)
          pacf_AIB = smt.pacf(prices['dAIB'], nlags=15)

          acf_BoI = smt.acf(prices['dBoI'], nlags=15)
          pacf_BoI = smt.pacf(prices['dBoI'], nlags=15)

          correlogram = pd.DataFrame({'acf_AIB':acf_AIB[1:], 'pacf_AIB':pacf_
          correlogram

          # 15 lags, the returns are daily so 15 days
```

```
/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/statioo
ls.py:667: FutureWarning:
```


fft=True will become the default after the release of the 0.12 release of statsmodels. To suppress this warning, explicitly set fft=False.

Out[107]:

	acf_AIB	pacf_AIB	acf_BOI	pacf_BOI
0	0.105898	0.106032	0.031277	0.031317
1	-0.028405	-0.040171	-0.030934	-0.032025
2	-0.058914	-0.052411	-0.040280	-0.038506
3	-0.090816	-0.081572	0.031261	0.033043
4	-0.036188	-0.022217	-0.016482	-0.021166
5	-0.053867	-0.057881	0.023390	0.025292
6	0.072869	0.075968	0.043883	0.044364
7	-0.020029	-0.051122	-0.025803	-0.030316
8	0.103871	0.111489	0.040641	0.049289
9	0.139978	0.117690	0.054519	0.052665
10	-0.018961	-0.033325	0.009914	0.004709
11	-0.033642	-0.015353	-0.023111	-0.013881
12	-0.012007	0.029166	0.020283	0.020901
13	0.066896	0.081205	0.070418	0.068023
14	-0.011839	-0.012036	0.021679	0.019680

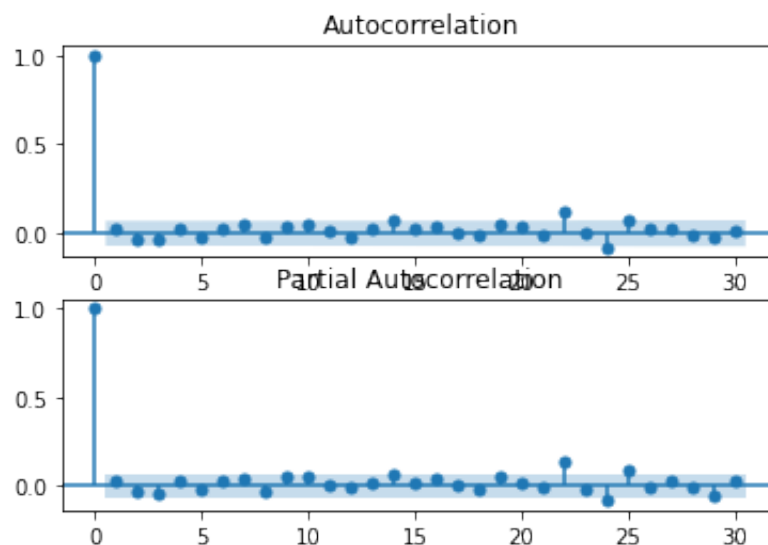

```
In [108]: # ACF and PACF of AIB Returns
from statsmodels.graphics.tsaplots import plot_acf
from statsmodels.graphics.tsaplots import plot_pacf

plt.figure()
plt.subplot(211)
plot_acf(prices['dAIB'], ax=plt.gca(), lags=30)
plt.subplot(212)
plot_pacf(prices['dAIB'], ax=plt.gca(), lags=30)
plt.show()
```



```
In [109]: # ACF and PACF of BoI Returns

plt.figure()
plt.subplot(211)
plot_acf(prices['dBoI'], ax=plt.gca(), lags=30)
plt.subplot(212)
plot_pacf(prices['dBoI'], ax=plt.gca(), lags=30)
plt.show()
```



Autocorrelation Figures

```
In [110]: autocorrelation_AIB = prices['dAIB'].autocorr()
print("The autocorrelation of daily AIB returns is %4.2f" %(autocor
```

The autocorrelation of daily AIB returns is 0.11

```
In [111]: autocorrelation_BoI = prices['dBoI'].autocorr()
print("The autocorrelation of daily BoI returns is %4.2f" %(autocor
```

The autocorrelation of daily BoI returns is 0.03

```
In [112]: num_obs = len(prices)
# Number of observations
```

```
In [113]: from math import sqrt

conf = 1.96/sqrt(num_obs)
print("The approximate confidence interval is +/- %4.2f" %(conf))
```

The approximate confidence interval is +/- 0.07

ADF Test

An Augmented Dicker Fuller (ADF) Test is a unit test for stationarity. The null hypothesis says $\alpha=1$ in the model equation.

As the null hypothesis assumes the presence of unit root ($\alpha=1$), the p-value should be less than 0.05 to reject the null hypothesis and infer that the series is stationary. A p-value > 0.05 would infer non-stationarity.

If a series has weak stationarity; the mean, variance and autocorrelation do not depend on time. If a series is not stationary, it will be difficult to model

```
In [114]: test1 = prices.loc['2020-02-05':]
train1 = prices.loc[:'2020-02-04']
```

```
In [115]: # Testing stationarity of AIB Stock Returns

def test_stationarity(timeseries):
    #Determining rolling statistics
    rolmean = prices['dAIB'].rolling(window=365).mean()
    rolstd = prices['dAIB'].rolling(window=365).std()


    #Plot rolling statistics:
    plt.plot(prices['dAIB'], color='blue', label='Original')
    plt.plot(rolmean, color='red', label='Rolling Mean')
    plt.plot(rolstd, color='black', label = 'Rolling Std')
```

```
plt.legend(loc='best')
plt.title('Rolling Mean and Standard Deviation')
plt.show(block=False)

print("Results of dickey fuller test")
adft = adfuller(timeseries,autolag='AIC')
# output for dft will give us without defining what the values are
#hence we manually write what values does it explains using a for
output = pd.Series(adft[0:4],index=['Test Statistics','p-value','No
for key,values in adft[4].items():
    output['critical value (%s)'%key] = values
print(output)

test_stationarity(train1['dAIB'])

# p-value < 0.05, therefore the time series is stationary
# Lack of a trend in the Rolling Mean and Stand Dev shows stationar.
```



```
Results of dickey fuller test
Test Statistics          -22.331378
p-value                 0.000000
No. of lags used        0.000000
Number of observations used  530.000000
critical value (1%)      -3.442749
critical value (5%)      -2.867009
critical value (10%)     -2.569683
dtype: float64
```

In [116]: # Testing stationarity of BoI Stock Returns

```
def test_stationarity(timeseries):
    #Determing rolling statistics
    rolmean = prices['dBoI'].rolling(window=365).mean()
    rolstd = prices['dBoI'].rolling(window=365).std()

    #Plot rolling statistics:
    plt.plot(prices['dBoI'], color='blue', label='Original')
    plt.plot(rolmean, color='red', label='Rolling Mean')
```

```


plt.plot(rolmean, color='red', label='Rolling Mean',
plt.plot(rolstd, color='black', label='Rolling Std')
plt.legend(loc='best')
plt.title('Rolling Mean and Standard Deviation')
plt.show(block=False)

print("Results of dickey fuller test")
adft = adfuller(timeseries,autolag='AIC')
# output for dft will give us without defining what the values are
#hence we manually write what values does it explains using a for
output = pd.Series(adft[0:4],index=['Test Statistics','p-value','No
for key,values in adft[4].items():
    output['critical value (%)'%key] = values
print(output)

test_stationarity(train1['dBoI'])

# p-value < 0.05, therefore the time series is stationary
# Lack of a trend in the Rolling Mean and Stand Dev shows stationar.

```



```

Results of dickey fuller test
Test Statistics          -23.683229
p-value                 0.000000
No. of lags used        0.000000
Number of observations used  530.000000
critical value (1%)      -3.442749
critical value (5%)      -2.867009
critical value (10%)     -2.569683
dtype: float64

```

```
In [117]: # Augmented Dickey Fuller Test for stationarity - AIB
from statsmodels.tsa.stattools import adfuller, kpss

result_AIB = adfuller(prices.dAIB.values, autolag='AIC')
print(f'ADF Statistic: {result_AIB[0]}')
print(f'p-value: {result_AIB[1]}')
for key, value in result_AIB[4].items():
    print('Critical Values:')
    print(f' {key},{value}')

# p-value is much less than significance level of 0.05, again confi
# p-vlaue less than 0.05, therefore we can reject the null hypotheses.
```

```
ADF Statistic: -7.558246463896617
p-value: 3.0585446331492924e-11
Critical Values:
 1%,-3.438783171038672
Critical Values:
 5%,-2.865262118650577
Critical Values:
10%,-2.568752018688748
```

```
In [118]: # Augmented Dickey Fuller Test for stationarity - BoI

result_BoI = adfuller(prices.dBoI.values, autolag='AIC')
print(f'ADF Statistic: {result_BoI[0]}')
print(f'p-value: {result_BoI[1]}')
for key, value in result_BoI[4].items():
    print('Critical Values:')
    print(f' {key},{value}')

# p-value is much less than significance level of 0.05, again confi
```

```
ADF Statistic: -27.164554340013876
p-value: 0.0
Critical Values:
 1%,-3.438686413400388
Critical Values:
 5%,-2.8652194721349424
Critical Values:
10%,-2.5687293001910008
```

KPSS Test

Kwiatkowski-Phillips-Schmidt-Shin is another test for stationarity. Key difference from the ADF test is that the null hypothesis here is that the series is actually stationary. Therefore, if the p-value < 0.05 for the KPSS test, the tested series is non-stationary.

```
In [119]: # KPSS Test for Stationarity - AIB returns
result_AIB2 = kpss(prices.dAIB.values, regression='c')
print('\nKPSS Statistic: %f' % result_AIB2[0])
print('p-value: %f' % result_AIB2[1])
for key, value in result_AIB2[3].items():
    print('Critical Values:')
    print(f' {key}: {value}')

# p-value > 0.05, so AIB returns are stationary
```

```
KPSS Statistic: 0.108159
p-value: 0.100000
Critical Values:
 10%: 0.347
Critical Values:
  5%: 0.463
Critical Values:
 2.5%: 0.574
Critical Values:
  1%: 0.739
```

```
/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/stattools.py:1875: FutureWarning:
```

The behavior of using `nlags=None` will change in release 0.13. Currently `nlags=None` is the same as `nlags="legacy"`, and so a sample-size lag length is used. After the next release, the default will change to be the same as `nlags="auto"` which uses an automatic lag length selection method. To silence this warning, either use `"auto"` or `"legacy"`

```
/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/stattools.py:1910: InterpolationWarning:
```

The test statistic is outside of the range of p-values available in the look-up table. The actual p-value is greater than the p-value returned.

```
In [120]: # KPSS Test for Stationarity - BoI returns
result_BoI2 = kpss(prices.dBoI.values, regression='c')
print('\nKPSS Statistic: %f' % result_BoI2[0])
print('p-value: %f' % result_BoI2[1])
for key, value in result_BoI2[3].items():
    print('Critical Values:')
    print(f' {key}: {value}')

# p-value > 0.05, so the BoI returns are stationary
```

KPSS Statistic: 0.156992

p-value: 0.100000

Critical Values:

10%: 0.347

Critical Values:

5%: 0.463

Critical Values:

2.5%: 0.574

Critical Values:

1%: 0.739

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/stattools.py:1910: InterpolationWarning:

The test statistic is outside of the range of p-values available in the look-up table. The actual p-value is greater than the p-value returned.

The ADF and KPSS tests have provided consistent results. Both tests show the returns are stationary.


Are the stocks cointegrated?

Cointegration is a long term correlation between two time series

In [121]: *# First plot the prices separately*


```
plt.subplot(2,1,1)
plt.plot(prices['AIB'], label='AIB', color='red')
plt.plot(prices['BoI'], label='BoI', color='blue')
plt.legend(loc='best', fontsize='small')
```

Out[121]: <matplotlib.legend.Legend at 0x7fac34316e20>



In [122]: *# Then plot the spread*
(The difference between the two series)

```
plt.subplot(2,1,2)
plt.plot(prices['BoI']-prices['AIB'], label='Spread')
plt.legend(loc='best', fontsize='small')
plt.axhline(y=0, linestyle='--', color='k')
plt.show()
```



In [123]: *# Reading CSV and transforming to returns again*

```
prices = pd.read_csv('IrishBanks.csv', index_col=0)
prices.index=pd.to_datetime(prices.index, infer_datetime_format=True)

# Transform stock prices to returns

prices['dAIB']= prices['AIB'].transform(lambda x : (x - x.shift(1)))
prices['dBoI']= prices['BoI'].transform(lambda x : (x - x.shift(1)))
prices = prices.dropna()
prices.head(3)
```

Out[123]:

	AIB	BoI	dAIB	dBoI
Date				
2018-01-03	5.435	7.370	-0.457875	2.432245
2018-01-04	5.450	7.545	0.275989	2.374491
2018-01-05	5.450	7.735	0.000000	2.518224

Choosing the Right Model

In [124]: `res_grid_AIB = smt.arma_order_select_ic(prices['dAIB'], max_ar=5, max_ma=5, fit_kw={'method':'css-mle', 'solver':'bfgs'})`

```
print("AIB's AIC")
print(res_grid_AIB.aic)
print("AIB's BIC")
print(res_grid_AIB.bic)

print(res_grid_AIB.aic_min_order)
print(res_grid_AIB.bic_min_order)

# AR on left of grid, MA on top
# Lowest BIC is (0,1)
```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/arma_model.py:472: FutureWarning:

statsmodels.tsa.arma_model.ARMA and statsmodels.tsa.arma_model.ARIMA have been deprecated in favor of statsmodels.tsa.arma.model.ARIMA (note the difference between arma and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arma.model.ARIMA makes use of the statespace framework and

is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```
import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARMA',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARIMA',
                        FutureWarning)
```

AIB's AIC

	0	1	2	3	4
0	4239.742145	4232.316699	4233.980140	4234.313432	4230.428997
1	4232.862823	4234.115510	4232.776366	4233.156728	4232.325456
2	4233.600092	4231.286373	4226.766873	4223.975109	4231.778120
3	4233.451308	4231.656318	4223.843649	4225.659227	4225.631624
4	4230.268716	4231.513625	4230.135319	4225.357985	4224.278309
5	4231.887147	4233.501176	4232.131124	4226.140746	4222.405513

	5
0	4232.376984
1	4234.311584
2	4233.739614
3	4226.004552
4	4217.088370
5	4222.240087

AIB's BIC

	0	1	2	3	4
0	4249.081141	4246.325193	4252.658132	4257.660922	4258.445986
1	4246.871317	4252.793502	4256.123857	4261.173716	4265.011943
2	4252.278084	4254.633863	4254.783861	4256.661595	4269.134104
3	4256.798799	4259.673307	4256.530136	4263.015212	4267.657107
4	4258.285704	4264.200112	4267.491303	4267.383468	4270.973290
5	4264.573634	4270.857161	4274.156607	4272.835726	4273.769992

	5
0	4265.063471
1	4271.667569
2	4275.765096
3	4272.699533
4	4268.452849
5	4278.274064

(4, 5)

(0, 1)

```
In [125]: res_arid_BoI = smt.arma_order_select_ic(prices['dBoI'].max_ar=5, m
```

```
fit_kw={'method':'css-mle', 'solver':'bfgs'})

print("BoI's AIC")
print(res_grid_BoI.aic)
print("BoI's BIC")
print(res_grid_BoI.bic)

print(res_grid_BoI.aic_min_order)
print(res_grid_BoI.bic_min_order)

# Lowest BIC is (0,0)
```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/arma_model.py:472: FutureWarning:

statsmodels.tsa.arma_model.ARMA and statsmodels.tsa.arma_model.ARIMA have been deprecated in favor of statsmodels.tsa.arma.model.ARIMA (note the difference between arma and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arma.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```
import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arma_model.ARMA',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arma_model.ARIMA',
                        FutureWarning)
```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/base/model.py:566: ConvergenceWarning:

Maximum Likelihood optimization failed to converge. Check mle_retvals

BoI's AIC	0	1	2	3	4
0	4117.299756	4118.479644	4119.887847	4120.457913	4121.249869
1	4118.528944	4120.325865	4121.471654	4121.372377	4123.001388
2	4119.726392	4121.330372	4122.889863	4118.576531	4120.830398

3	4120.566571	4121.664863	4119.295036	4120.344456	4122.042012
4	4121.717677	4123.420077	4113.793553	4120.818948	4122.202824
5	4123.371347	4125.284915	4117.744296	4115.384376	4119.048339

	5
0	4122.964717
1	4121.143434
2	4122.799270
3	4114.636237
4	4120.088931
5	4120.893646

BoI's BIC

	0	1	2	3	4
\					
0	4126.638752	4132.488138	4138.565840	4143.805404	4149.266858
1	4132.537438	4139.003857	4144.819144	4149.389366	4155.687875
2	4138.404384	4144.677863	4150.906851	4151.263018	4158.186382
3	4143.914062	4149.681851	4151.981523	4157.700441	4164.067494
4	4149.734666	4156.106564	4151.149538	4162.844431	4168.897805
5	4156.057833	4162.640900	4159.769779	4162.079357	4170.412818

	5
0	4155.651203
1	4158.499419
2	4164.824753
3	4161.331218
4	4171.453410
5	4176.927623

(4, 2)
(0, 0)

Part 2(a) Forecasting using ARMA Processes

2A: Show how you might forecast both series individually using their ARMA processes


```
In [126]: # Examine the seasonality of the returns
from statsmodels.tsa.seasonal import seasonal_decompose
decomposition = seasonal_decompose(prices.dAIB, freq = 260) #Was the
trend = decomposition.trend
seasonal = decomposition.seasonal
residual = decomposition.resid

plt.figure(figsize=(12,6))
plt.subplot(411)
plt.plot(prices['dAIB'], label='Original')
plt.legend(loc='best')
plt.subplot(412)
plt.plot(trend, label='Trend')
plt.legend(loc='best')
plt.subplot(413)
plt.plot(seasonal, label='Seasonality')
plt.legend(loc='best')
plt.subplot(414)
plt.plot(residual, label='Residuals')
plt.legend(loc='best')
plt.tight_layout()
plt.show()

# Shows stationarity in the returns, no trend
```

/var/folders/mh/j7bgd0wx1mb1jjnzcp0c8dfw0000gn/T/ipykernel_83727/29824519.py:3: FutureWarning:

the 'freq' keyword is deprecated, use 'period' instead




```
In [127]: # Examine the seasonality
decomposition = seasonal_decompose(prices.dBoI, freq = 260) #Was th
trend = decomposition.trend
seasonal = decomposition.seasonal
residual = decomposition.resid

plt.figure(figsize=(12,6))
plt.subplot(411)
plt.plot(prices['dBoI'], label='Original')
plt.legend(loc='best')
plt.subplot(412)
plt.plot(trend, label='Trend')
plt.legend(loc='best')
plt.subplot(413)
plt.plot(seasonal, label='Seasonality')
plt.legend(loc='best')
plt.subplot(414)
plt.plot(residual, label='Residuals')
plt.legend(loc='best')
plt.tight_layout()
plt.show()

# Shows stationarity in the returns, no trend
```

/var/folders/mh/j7bgd0wx1mb1jjnzcp0c8dfw0000gn/T/ipykernel_83727/200272331.py:2: FutureWarning:

the 'freq' keyword is deprecated, use 'period' instead



```
In [128]: # Make an Auto Regressive Moving Average model for AIB returns
# This function finds the best parameters for Ar(p) and Ma(q)
from statsmodels.tsa.stattools import ARMA
def best_AR_MA_checker(df, lower, upper):
    from statsmodels.tsa.stattools import ARMA
    from statsmodels.tsa.stattools import adfuller
    arg=np.arange(lower,upper)
    arg1=np.arange(lower,upper)
    best_param_i=0
    best_param_j=0
    temp=12000000
    rs=99
    for i in arg:
        for j in arg1:
            model=ARMA(df, order=(i,0,j))
            result=model.fit(disp=0)
            resid=adfuller(result.resid)
            if (result.aic<temp and adfuller(result.resid)[1]<0.05):
                temp=result.aic
                best_param_i=i
                best_param_j=j
                rs=resid[1]

    print ("AR: %d, MA: %d, AIC: %d; resid stationarity che
    print("the following function prints AIC criteria and finds the
    print("best AR: %d, best MA: %d, best AIC: %d; resid stationar
best_AR_MA_checker(prices.dAIB,0,4) #For each parameter I want to t
```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/arma_model.py:472: FutureWarning:

statsmodels.tsa.arma_model.ARMA and statsmodels.tsa.arma_model.ARIMA have been deprecated in favor of statsmodels.tsa.arma.model.ARIMA (note the . between arma and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arma.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```
In [129]: # Above output might be difficult to read
# It states best ARMA Model is (1,0) for AIB
```

```

In [130]: # Make an Auto Regressive Moving Average model for BoI returns
# This code finds the best parameters for Ar(p) and Ma(q)
def best_AR_MA_checker(df, lower, upper):
    from statsmodels.tsa.stattools import ARMA
    from statsmodels.tsa.stattools import adfuller
    arg=np.arange(lower,upper)
    arg1=np.arange(lower,upper)
    best_param_i=0
    best_param_j=0
    temp=12000000
    rs=99
    for i in arg:
        for j in arg1:
            model=ARMA(df, order=(i,0,j))
            result=model.fit(dis=0)
            resid=adfuller(result.resid)
            if (result.aic<temp and adfuller(result.resid)[1]<0.05):
                temp=result.aic
                best_param_i=i
                best_param_j=j
                rs=resid[1]

    print ("AR: %d, MA: %d, AIC: %d; resid stationarity che

    print("the following function prints AIC criteria and finds the
    print("best AR: %d, best MA: %d, best AIC: %d; resid stationar
best_AR_MA_checker(prices.dBoI,0,4) #For each parameter I want to t

```

RIMA have

been deprecated in favor of statsmodels.tsa.arima.model.ARIMA (not
e the .
between arima and model) and
statsmodels.tsa.SARIMAX. These will be removed after the 0.12 rele
ase.

statsmodels.tsa.arima.model.ARIMA makes use of the statespace fram
ework and
is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until th
ey are
removed, use:

```

import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARM
A',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARI

```

```

In [131]: # Above says best ARMA model for BoI is (0,0)
# ARMA (1,0) has a very similar AIC

```

```

In [132]: from statsmodels.tsa.stattools import acf, pacf

```



```

In [132]: from statsmodels.tsa.stattools import acf, pacf

lag_acf = acf(prices['dAIB'], nlags=5)
lag_pacf = pacf(prices['dAIB'], nlags=5, method='ols')

# Plot ACF
plt.figure(figsize=(15,5))
plt.subplot(121)
plt.plot(lag_acf)
plt.axhline(y=0, linestyle='--', color='gray')
plt.axhline(y=-1.96/np.sqrt(len(prices['dAIB'])), linestyle='--', color='gray')
plt.axhline(y=1.96/np.sqrt(len(prices['dAIB'])), linestyle='--', color='gray')
plt.title('AIB Autocorrelation Function')

# Plot PACF
plt.figure(figsize=(15,5))
plt.subplot(122)
plt.plot(lag_pacf)
plt.axhline(y=0, linestyle='--', color='gray')
plt.axhline(y=-1.96/np.sqrt(len(prices['dAIB'])), linestyle='--', color='gray')
plt.axhline(y=1.96/np.sqrt(len(prices['dAIB'])), linestyle='--', color='gray')
plt.title('AIB Partial Autocorrelation Function')

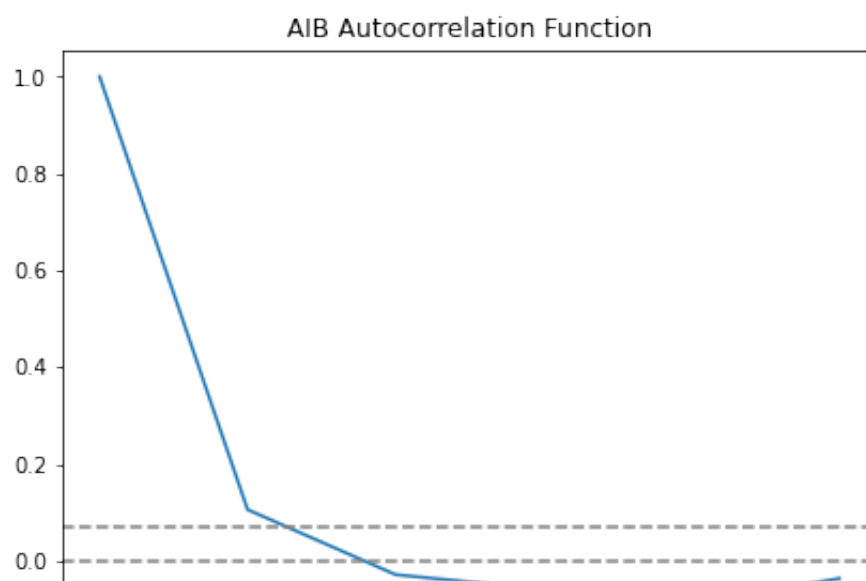
# Look to see where the graph hits zero for the first time
# It cuts it at approximately 2 on Autocorrelation Graph
# Approx 2 on the partial autocorrelation graph
# Therefore the ARMA (p,q) model for AIB might be:
# ARMA (2,2)


```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/stattools.py:667: FutureWarning:

fft=True will become the default after the release of the 0.12 release of statsmodels. To suppress this warning, explicitly set fft=False.

Out[132]: Text(0.5, 1.0, 'AIB Partial Autocorrelation Function')





```
In [133]: lag_acf = acf(prices['dBoI'], nlags=5)
lag_pacf = pacf(prices['dBoI'], nlags=5, method='ols')

# Plot ACF
plt.figure(figsize=(15,5))
plt.subplot(121)
plt.plot(lag_acf)
plt.axhline(y=0, linestyle='--', color='gray')
plt.axhline(y=-1.96/np.sqrt(len(prices['dBoI'])), linestyle='--', color='gray')
plt.axhline(y=1.96/np.sqrt(len(prices['dBoI'])), linestyle='--', color='gray')
plt.title('BoI Autocorrelation Function')


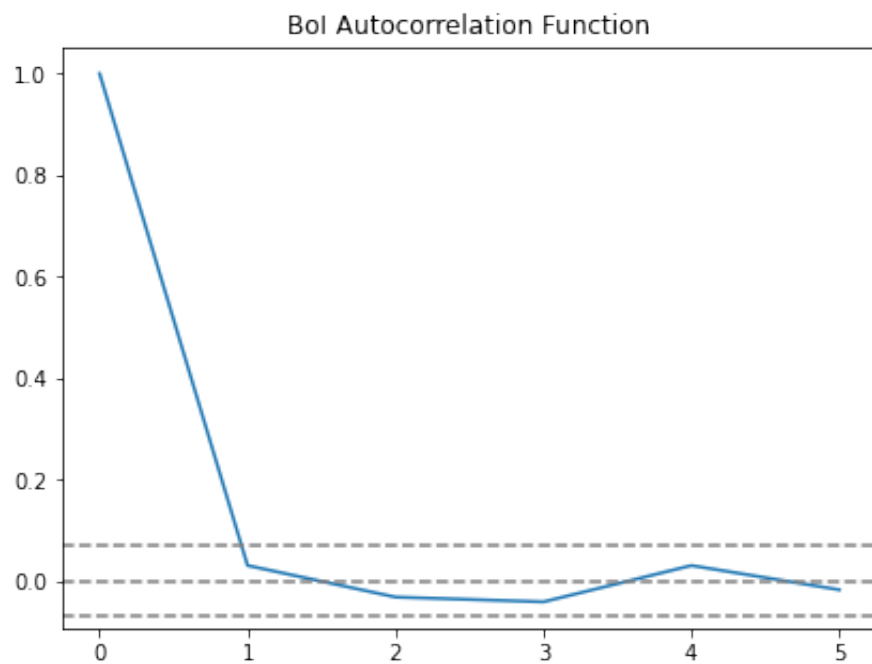
# Plot PACF
plt.figure(figsize=(15,5))
plt.subplot(122)
plt.plot(lag_pacf)
plt.axhline(y=0, linestyle='--', color='gray')
plt.axhline(y=-1.96/np.sqrt(len(prices['dBoI'])), linestyle='--', color='gray')
plt.axhline(y=1.96/np.sqrt(len(prices['dBoI'])), linestyle='--', color='gray')
plt.title('BoI Partial Autocorrelation Function')

# Look to see where the graph hits zero for the first time
# It cuts it at approximately 2 on Autocorrelation Graph
# Approx 1 on the partial autocorrelation graph
# Therefore the ARMA (p,q) model for BoI might be:
# ARMA (2,1)
```

```
/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/stattools.py:667: FutureWarning:
```

`fft=True` will become the default after the release of the 0.12 release of statsmodels. To suppress this warning, explicitly set `fft=False`.


Out[133]: Text(0.5, 1.0, 'BoI Partial Autocorrelation Function')



```
In [134]: from statsmodels.graphics.tsaplots import plot_acf
from statsmodels.graphics.tsaplots import plot_pacf


# Compare ACF and PACF for AIB returns
fig, axes = plt.subplots(2,1)

# Plot ACF
plot_acf(prices['dAIB'], lags = 15, ax=axes[0])
#Plot PACF
plot_pacf(prices['dAIB'], lags=15, ax=axes[1])
plt.show()
```



```
In [135]: # Compare ACF and PACF for BoI returns
fig, axes = plt.subplots(2,1)

# Plot ACF
plot_acf(prices['dBoI'], lags = 15, ax=axes[0])
#Plot PACF
plot_pacf(prices['dBoI'], lags=15, ax=axes[1])
plt.show()
```



```
In [136]: # Estimate the AR for AIB returns
# Fit the data to an AR(p) for p = 0,...,6 , and save the BIC

from statsmodels.tsa.arima_model import ARMA

BIC = np.zeros(7)
for p in range(7):
    mod = ARMA(prices['dAIB'], order=(p,0))
    res = mod.fit()
# Save BIC for AR(p)
    BIC[p] = res.bic

# Plot the BIC as a function of p
plt.figure(figsize=(10,5))
plt.plot(range(1,7), BIC[1:7], marker='o')
plt.xlabel('Order of AR Model for AIB Returns')
plt.ylabel('Bayesian Information Criterion')
plt.show()

# Minimum AR at 1
```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/arima_model.py:472: FutureWarning:

statsmodels.tsa.arima_model.ARMA and statsmodels.tsa.arima_model.ARIMA have been deprecated in favor of statsmodels.tsa.arima.model.ARIMA (note the . between arima and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arima.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```
In [137]: from statsmodels.tsa.arima_model import ARMA
# Forecast AIB returns using the first AR(1) model

mod = ARMA(prices['dAIB'], order=(1,0))
res = mod.fit()
res.plot_predict(start='2020-11-05', end='2021-02-05')
plt.show()

# Arma model (1,0) used to predict last 6 months prices
```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/arima_model.py:472: FutureWarning:

statsmodels.tsa.arima_model.ARMA and statsmodels.tsa.arima_model.ARIMA have been deprecated in favor of statsmodels.tsa.arima.model.ARIMA (note the . between arima and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arima.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```
import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARMA',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARIMA',
                        FutureWarning)
```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/tsa_model.py:581: ValueWarning:

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

This problem is unconstrained.

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 12

At X0 0 variables are exactly at the bounds

At iterate 0 f= 2.68202D+00 |proj g|= 6.64357D-05

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient


F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
2	4	6	1	0	0	0.000D+00	2.682D+00

F = 2.6820195575481245

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL




```
In [138]: # Estimate the AR for BoI returns
# Fit the data to an AR(p) for p = 0,...,6 , and save the BIC
BIC = np.zeros(7)
for p in range(7):
    mod = ARMA(prices['dBoI'], order=(p,0))
    res = mod.fit()
# Save BIC for AR(p)
BIC[p] = res.bic

# Plot the BIC as a function of p
plt.figure(figsize=(10,5))
plt.plot(range(1,7), BIC[1:7], marker='o')
plt.xlabel('Order of AR Model for BoI Returns')
plt.ylabel('Bayesian Information Criterion')
plt.show()

# Minimum AR at 1
```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/arma_model.py:472: FutureWarning:

statsmodels.tsa.arma_model.ARMA and statsmodels.tsa.arma_model.ARIMA have been deprecated in favor of statsmodels.tsa.arma.model.ARIMA (note the . between arma and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arma.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```
In [139]: # Forecast BoI returns using the first AR(1) model

mod = ARMA(prices['dBoI'], order=(1,0))
res = mod.fit()
res.plot_predict(start='2020-11-05', end='2021-02-05')
plt.show()

# Arma model (1,0) used to predict last 6 months prices
```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/arma_model.py:472: FutureWarning:

statsmodels.tsa.arma_model.ARMA and statsmodels.tsa.arma_model.ARIMA have

been deprecated in favor of `statsmodels.tsa.arima.model.ARIMA` (note the . between `arima` and `model`) and `statsmodels.tsa.SARIMAX`. These will be removed after the 0.12 release.

`statsmodels.tsa.arima.model.ARIMA` makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```
import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARMA',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARIMA',
                        FutureWarning)
```

`/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/tsa_model.py:581: ValueWarning:`

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

This problem is unconstrained.

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 12

At X0 0 variables are exactly at the bounds

At iterate 0 f= 2.60947D+00 |proj g|= 2.90878D-04

At iterate 5 f= 2.60947D+00 |proj g|= 0.00000D+00

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient


F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
2	5	8	1	0	0	0.000D+00	2.609D+00

F = 2.6094726799406538

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL



MA Models

```
In [140]: # Forecast AIB returns using the first MA(1) model

mod = ARMA(prices['dAIB'], order=(0,1))
res = mod.fit()
res.plot_predict(start = '2020-11-05' ,end = '2021-02-05' )
plt.show()
```

Arma model (0,1) used to predict last 6 months prices

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/arma_model.py:472: FutureWarning:

statsmodels.tsa.arma_model.ARMA and statsmodels.tsa.arma_model.ARIMA have been deprecated in favor of statsmodels.tsa.arma.model.ARIMA (note the . between arma and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arma.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until the

ey are
removed, use:

```
import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARM
A',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARI
MA',
                        FutureWarning)
```

```
/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/ts
a_model.py:581: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

This problem is unconstrained.

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 12

At X0 0 variables are exactly at the bounds

At iterate 0 f= 2.68167D+00 |proj g|= 6.99441D-05

At iterate 5 f= 2.68167D+00 |proj g|= 0.00000D+00

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient


F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
2	5	8	1	0	0	0.000D+00	2.682D+00
F =	2.6816730321496633						

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL





In [141]: *# Forecast BoI returns using the first MA(1) model*

```
mod = ARMA(prices['dBoI'], order=(0,1))
res = mod.fit()
res.plot_predict(start = '2020-11-05' ,end = '2021-02-05' )
plt.show()
```

Arma model (0,1) used to predict last 6 months prices

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/arma_model.py:472: FutureWarning:

statsmodels.tsa.arma_model.ARMA and statsmodels.tsa.arma_model.ARIMA have been deprecated in favor of statsmodels.tsa.arma.model.ARIMA (note the difference between arma and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arma.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```
import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arma_model.ARMA',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arma_model.ARIMA',
                        FutureWarning)
```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/tsa_model.py:581: ValueWarning:

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

This problem is unconstrained.

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 12

At X0 0 variables are exactly at the bounds

At iterate 0 f= 2.60944D+00 |proj g|= 1.01696D-05

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point


Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
2	4	6	1	0	0	0.000D+00	2.609D+00
F = 2.6094413981716196							

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL



But Which ARMA Model is Best?

In [142]: *# Find Best ARMA Model for AIB*

```
# Fit the data to an AR(1) model and print AIC:
mod_ar1 = ARMA(prices['dAIB'], order=(1, 0))
res_ar1 = mod_ar1.fit()
print("The AIC for an AR(1) is: ", res_ar1.aic)

# Fit the data to an AR(2) model and print AIC:
mod_ar2 = ARMA(prices['dAIB'], order=(2, 0))
res_ar2 = mod_ar2.fit()
print("The AIC for an AR(2) is: ", res_ar2.aic)

# Fit the data to an ARMA(1,1) model and print AIC:
mod_arma11 = ARMA(prices['dAIB'], order = (1,1))
res_arma11 = mod_arma11.fit()
print("The AIC for an ARMA(1,1) is: ", res_arma11.aic)

# AR(1) has the lowest AIC score of the three models – for AIB
```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/arma_model.py:472: FutureWarning:

statsmodels.tsa.arma_model.ARMA and statsmodels.tsa.arma_model.ARIMA have been deprecated in favor of statsmodels.tsa.arma.model.ARIMA (note the difference between arma and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arma.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

In [143]: *# Find Best ARMA Model for BoI*

```
# Fit the data to an AR(1) model and print AIC:
mod_ar1 = ARMA(prices['dBoI'], order=(1, 0))
res_ar1 = mod_ar1.fit()
print("The AIC for an AR(1) is: ", res_ar1.aic)

# Fit the data to an AR(2) model and print AIC:
mod_ar2 = ARMA(prices['dBoI'], order=(2, 0))
res_ar2 = mod_ar2.fit()
print("The AIC for an AR(2) is: ", res_ar2.aic)

# Fit the data to an ARMA(1,1) model and print AIC:
mod_arma11 = ARMA(prices['dBoI'], order = (2,1))
res_arma11 = mod_arma11.fit()
print("The AIC for an ARMA(2,1) is: ", res_arma11.aic)

# Again AR(1) has the lowest AIC score of the three models for BoI
```

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 12

At X0 0 variables are exactly at the bounds

At iterate 0 f= 2.60947D+00 |proj g|= 2.90878D-04

At iterate 5 f= 2.60947D+00 |proj g|= 0.00000D+00

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Next = number of active bounds at first nonoptimal Cauchy point

In [144]: *# Forecast AIB Returns using an ARIMA(1,0,0) model*

```
from statsmodels.tsa.arima_model import ARIMA

mod = ARIMA(prices['dAIB'], order=(1,0,0))
res = mod.fit()

# Plot the original series and the forecasted series
res.plot_predict(start='2020-08-05', end='2021-02-05')
plt.show()
```

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 12

At X0 0 variables are exactly at the bounds

At iterate 0 f= 2.68202D+00 |proj g|= 6.64357D-05

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *


N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
2	4	6	1	0	0	0.000D+00	2.682D+00
F = 2.6820195575481245							

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/tsa_model.py:581: ValueWarning:

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

This problem is unconstrained.



In [145]: # Forecast BoI Returns using an ARIMA(1,0,0) model

```
mod = ARIMA(prices['dBoI'], order=(1,0,0))
res = mod.fit()
```

```
# Plot the original series and the forecasted series
res.plot_predict(start='2020-08-05', end='2021-02-05')
plt.show()
```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/arma_model.py:472: FutureWarning:

statsmodels.tsa.arma_model.ARMA and statsmodels.tsa.arma_model.ARIMA have been deprecated in favor of statsmodels.tsa.arma.model.ARIMA (note the difference between arma and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arma.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```
import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arma_model.ARMA',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arma_model.ARIMA',
                        FutureWarning)
```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/tsa_model.py:581: ValueWarning:

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

This problem is unconstrained.

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 12

At X0 0 variables are exactly at the bounds

At iterate 0 f= 2.60947D+00 |proj g|= 2.90878D-04

At iterate 5 f= 2.60947D+00 |proj g|= 0.00000D+00

* * *


Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
2	5	8	1	0	0	0.000D+00	2.609D+00

F = 2.6094726799406538

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL



In [146]: # Model AIB using BIC

```

model_AIB = smt.ARIMA(prices['dAIB'], order=(1,0,0))
res_AIB = model_AIB.fit()
print(res_AIB.summary())

```

```

# No differencing needs to be completed as the returns are stationa
# 1 Lag of AR, 0 Lags of MA
# Lag 1 is significant, p-value < 0.003

```

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 12

At X0 0 variables are exactly at the bounds

At iterate 0 f= 2.68202D+00 |proj g|= 6.64357D-05

* * *

Tit = total number of iterations
 Tnf = total number of function evaluations
 Tnint = total number of segments explored during Cauchy searches
 Skip = number of BFGS updates skipped
 Nact = number of active bounds at final generalized Cauchy point
 Projg = norm of the final projected gradient
 F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
2	4	6	1	0	0	0.000D+00	2.682D+00

F = 2.6820195575481245

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

ARMA Model Results

```

=====
=====
Dep. Variable:          dAIB    No. Observations:
788
Model:                ARMA(1, 0)    Log Likelihood
-2113.431
Method:                css-mle    S.D. of innovations
3.536
Date:                  Fri, 19 Nov 2021    AIC
4232.863
Time:                  10:46:32    BIC
4246.871
Sample:                0    HQIC
4238.248
  
```

```

=====
=====
              coef      std err          z      P>|z|      [0.025
0.975]
-----
const          -0.0934      0.141      -0.663      0.507      -0.369
0.183
ar.L1.dAIB       0.1058      0.035       2.988      0.003       0.036
0.175
  
```

Roots

```

=====
=====
              Real          Imaginary      Modulus
Frequency
-----
AR.1          9.4510          +0.0000j      9.4510
0.0000
  
```

```
-----
/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/arma_model.py:472: FutureWarning:
```

statsmodels.tsa.arma_model.ARMA and statsmodels.tsa.arma_model.ARIMA have been deprecated in favor of statsmodels.tsa.arma.model.ARIMA (note the difference between arma and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arma.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```
import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arma_model.ARMA',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arma_model.ARIMA',
                        FutureWarning)
```

```
/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/tsa_model.py:581: ValueWarning:
```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

This problem is unconstrained.

In [147]: *# Model BoI using BIC*

```
model_BoI = smt.ARIMA(prices['dBoI'], order=(1,0,0))
res_BoI = model_BoI.fit()
print(res_BoI.summary())
```

```
# No differencing needs to be completed as the returns are stationary
# 1 Lag of AR, 0 Lags of MA
```

```
/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/arma_model.py:472: FutureWarning:
```

statsmodels.tsa.arima_model.ARMA and statsmodels.tsa.arima_model.ARIMA have been deprecated in favor of statsmodels.tsa.arima.model.ARIMA (note the difference between arima and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arima.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```
import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARMA',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARIMA',
                        FutureWarning)
```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/tsa_model.py:581: ValueWarning:

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

This problem is unconstrained.

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 12

At X0 0 variables are exactly at the bounds

At iterate 0 f= 2.60947D+00 |proj g|= 2.90878D-04

At iterate 5 f= 2.60947D+00 |proj g|= 0.00000D+00

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
2	5	8	1	0	0	0.000D+00	2.609D+00

F = 2.6094726799406538

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

ARMA Model Results

=====

=====

Dep. Variable: dBoI No. Observations:

788

Model: ARMA(1, 0) Log Likelihood

-2056.264

Method: css-mle S.D. of innovations

3.289

Date: Fri, 19 Nov 2021 AIC

4118.529

Time: 10:46:32 BIC

4132.537

Sample: 0 HQIC

4123.914

=====

=====

	coef	std err	z	P> z	[0.025
--	------	---------	---	------	--------

0.975]

const	-0.0453	0.121	-0.374	0.708	-0.282
-------	---------	-------	--------	-------	--------

0.192

ar.L1.dBoI	0.0313	0.036	0.878	0.380	-0.039
------------	--------	-------	-------	-------	--------

0.101

Roots

=====

=====

	Real	Imaginary	Modulus
--	------	-----------	---------

Frequency

AR.1	31.9898	+0.0000j	31.9898
------	---------	----------	---------

0.0000

```
In [148]: # Two types of useful forecasts
# - static
# - dynamic

# Dynamic forecasts calculate multi-step forecasts starting from the
# Static forecasts imply a sequence of one-step ahead forecasts, rol
```

```
In [149]: # Static - AIB

model_all_AIB = smt.ARIMA(prices['dAIB'], order =(1,0,0))
res_all_AIB = model_all_AIB.fit()
res_all_AIB.plot_predict('2020-08-05', '2021-02-05', dynamic=False)

# Last 6 months of prices
# Visually, Static looks like a better model than Dynamic for AIB r
```

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 12

At X0 0 variables are exactly at the bounds

At iterate 0 f= 2.68202D+00 |proj g|= 6.64357D-05

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
2	4	6	1	0	0	0.000D+00	2.682D+00
F =	2.6820195575481245						

CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/arma_model.py:472: FutureWarning:

statsmodels.tsa.arma_model.ARMA and statsmodels.tsa.arma_model.ARIMA have
been deprecated in favor of statsmodels.tsa.arma.model.ARIMA (not
e the .

between arima and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arima.model.ARIMA makes use of the statespace framework and is both well tested and maintained.


To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```
import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARMA',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARIMA',
                        FutureWarning)
```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/tsa_model.py:581: ValueWarning:

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

This problem is unconstrained.



In [150]: *# Dynamic - AIB*

```
res_all_AIB.plot_predict('2020-08-05', '2021-02-05', dynamic=True);  
# Dynamic is much flatter
```



In [151]: *# Static - BoI*

```
model_all_BoI = smt.ARIMA(prices['dBoI'], order=(1,0,0))  
res_all_BoI = model_all_BoI.fit()  
res_all_BoI.plot_predict('2020-08-05', '2021-02-05', dynamic=False)  
# Last 6 months of prices
```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/arma_model.py:472: FutureWarning:

statsmodels.tsa.arma_model.ARMA and statsmodels.tsa.arma_model.ARIMA have been deprecated in favor of statsmodels.tsa.arma.model.ARIMA (note the . between arma and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arma.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```
import warnings  
warnings.filterwarnings('ignore', 'statsmodels.tsa.arma_model.ARMA',
```

```

FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arima_model.ARI
MA',
FutureWarning)

```

```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/ts
a_model.py:581: ValueWarning:

```

A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

This problem is unconstrained.

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 2 M = 12

At X0 0 variables are exactly at the bounds

At iterate 0 f= 2.60947D+00 |proj g|= 2.90878D-04

At iterate 5 f= 2.60947D+00 |proj g|= 0.00000D+00

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point


Projg = norm of the final projected gradient


F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
2	5	8	1	0	0	0.000D+00	2.609D+00
F =	2.6094726799406538						


CONVERGENCE: NORM_OF_PROJECTED_GRADIENT_<=_PGTOL





In [152]: *# Dynamic - BoI*

```
res_all_BoI.plot_predict('2020-08-05', '2021-02-05', dynamic=True);
```



In [153]: *# Check the accuracy of the forecast – AIB*

```
def rmse(pred, target):
    return np.sqrt(((pred - target)**2).mean())

prices_outsample = prices['2019-02-05': '2021-02-05']

pred1= res_all_AIB.predict('2019-02-05', '2021-02-05', dynamic=False)
stats1= rmse(pred1, prices_outsample['dAIB'])
print('root mean squared error, static: {}'.format(stats1))

pred2= res_all_AIB.predict('2019-02-05', '2021-02-05', dynamic=True)
stats2= rmse(pred2, prices_outsample['dAIB'])
print('root mean squared error, dynamic: {}'.format(stats2))

prices_outsample2 = prices['2018-01-03': '2021-02-05']
pred3= res_all_AIB.predict('2018-01-03', '2021-02-05', dynamic=False)
stats3= rmse(pred3, prices_outsample2['dAIB'])
print('root mean squared error, long (whole period): {}'.format(stats3))

# The closer to 0 the better
# The static model is a slightly better predictor than the dynamic
# The static graphs did look better than the dynamic
```

```
root mean squared error, static: 4.208232505366341
root mean squared error, dynamic: 4.239276415649806
root mean squared error, long (whole period): 3.536275151420149
```

In [154]: *# Check the accuracy of the forecast - BoI*

```
def rmse(pred, target):
    return np.sqrt(((pred - target)**2).mean())

prices_outsample = prices['2019-02-05': '2021-02-05']

pred4= res_all_BoI.predict('2019-02-05', '2021-02-05', dynamic=False)
stats4= rmse(pred4, prices_outsample['dBoI'])
print('root mean squared error, static: {}'.format(stats4))

pred5= res_all_BoI.predict('2019-02-05', '2021-02-05', dynamic=True)
stats5= rmse(pred5, prices_outsample['dBoI'])
print('root mean squared error, dynamic: {}'.format(stats5))

prices_outsample2 = prices['2018-01-03': '2021-02-05']
pred6= res_all_BoI.predict('2018-01-03', '2021-02-05', dynamic=False)
stats6= rmse(pred3, prices_outsample2['dBoI'])
print('root mean squared error, long (whole period): {}'.format(stats6))

# The closer to 0 the better
# The static and dynamic models are similar for BoI

root mean squared error, static: 3.845590536643706
root mean squared error, dynamic: 3.849923383394704
root mean squared error, long (whole period): 3.278440574670801
```

Another ARIMA Method: Forecasting prices

In [155]: **from** sklearn.metrics **import** mean_squared_error

```
# Split AIB Prices into train and test data
train_data, test_data = prices[0:int(len(prices)*0.7)], prices[int(len(prices)*0.7):]

training_data = train_data['AIB'].values
test_data = test_data['AIB'].values

history = [x for x in training_data]
model_predictions = []
N_test_observations = len(test_data)

# ARIMA (1,1,0) used
# As it is stock prices and not returns, 1 order of differencing needed
for time_point in range(N_test_observations):
    model = ARIMA(history, order=(1,1,0))
    model_fit = model.fit(disp=0)
    output = model_fit.forecast()
    yhat = output[0]
    model_predictions.append(yhat)
    true_test_value = test_data[time_point]
    history.append(true_test_value)
```

```
MSE_error = mean_squared_error(test_data, model_predictions)
print('Testing Mean Squared Error is {}'.format(MSE_error))
```

```
# Mean Squared Error is low
```

```
/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/arma_m
odel.py:472: FutureWarning:
```

statsmodels.tsa.arma_model.ARMA and statsmodels.tsa.arma_model.A
RIMA have
been deprecated in favor of statsmodels.tsa.arma.model.ARIMA (not
e the .
between arma and model) and
statsmodels.tsa.SARIMAX. These will be removed after the 0.12 rele
ase.

statsmodels.tsa.arma.model.ARIMA makes use of the statespace fram
ework and
is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until th
ey are
removed, use:

```
import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arma_model.ARM
A',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arma_model.ARI
MA',
                        FutureWarning)
```


Testing Mean Squared Error is 0.004069872225705441

```
In [156]: prices['date'] = prices.index
```

```
In [157]: test_set_range = prices[int(len(prices)*0.7):].index

plt.figure(figsize=(13,8))
plt.plot(test_set_range, model_predictions, color='blue',
         marker='o', linestyle='dashed', label='Predicted Price')
plt.plot(test_set_range, test_data, color='red', label='Actual Price')
plt.title('AIB Stock Price Prediction')
plt.xlabel('Date')
plt.ylabel('Prices')
plt.legend()
plt.show()

# Predicted price maps the actual very well
```



Now for BoI Stock Price

```
In [158]: # Split BoI Prices into train and test data
train_data, test_data = prices[0:int(len(prices)*0.7)], prices[int(

training_data = train_data['BoI'].values
test_data = test_data['BoI'].values

history = [x for x in training_data]
model_predictions = []
N_test_observations = len(test_data)

# ARIMA (1,1,0) used
# As it is stock prices and not returns, 1 order of differencing ne
for time_point in range(N_test_observations):
    model = ARIMA(history, order=(1,1,0))
```



```

model_fit = model.fit(dispatch=0)
output = model_fit.forecast()
yhat = output[0]
model_predictions.append(yhat)
true_test_value = test_data[time_point]
history.append(true_test_value)

MSE_error = mean_squared_error(test_data, model_predictions)
print('Testing Mean Squared Error is {}'.format(MSE_error))

# Mean squared error again is low

```

```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/arma_model.py:472: FutureWarning:

```

statsmodels.tsa.arma_model.ARMA and statsmodels.tsa.arma_model.ARIMA have been deprecated in favor of statsmodels.tsa.arma.model.ARIMA (note the . between arma and model) and statsmodels.tsa.SARIMAX. These will be removed after the 0.12 release.

statsmodels.tsa.arma.model.ARIMA makes use of the statespace framework and is both well tested and maintained.

To silence this warning and continue using ARMA and ARIMA until they are removed, use:

```

import warnings
warnings.filterwarnings('ignore', 'statsmodels.tsa.arma_model.ARMA',
                        FutureWarning)
warnings.filterwarnings('ignore', 'statsmodels.tsa.arma_model.ARIMA',
                        FutureWarning)

```

```

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/base/model.py:566: ConvergenceWarning:

```

Maximum Likelihood optimization failed to converge. Check mle_retvals

Testing Mean Squared Error is 0.010094015345757164


```
In [159]: test_set_range = prices[int(len(prices)*0.7):].index

plt.figure(figsize=(13,8))
plt.plot(test_set_range, model_predictions, color='blue',
         marker='o', linestyle='dashed', label='Predicted Price')

plt.plot(test_set_range, test_data, color='red', label='Actual Price')

plt.title('BoI Stock Price Prediction')
plt.xlabel('Date')
plt.ylabel('Prices')
plt.legend()
plt.show()

# Again, predicted price maps the actual very well
```



Part 2(b) Analyse Volatility using GARCH

2B: Analyse the volatility of both series individually and demonstrate how a GARCH process can be used to model their volatility

```
In [160]: from statsmodels.stats.diagnostic import het_arch
from statsmodels.compat import lzip
```

In [161]: *# Use a specialist module for ARCH*

```
import sys
!{sys.executable} -m pip install arch
```

Collecting arch

Downloading arch-5.0.1-cp39-cp39-macosx_10_9_x86_64.whl (876 kB)
 |████████████████████| 876 kB 1.9 MB/s eta 0:00:01

1
 Collecting property-cached<=1.6.4

Downloading property_cached-1.6.4-py2.py3-none-any.whl (7.8 kB)
 Requirement already satisfied: numpy>=1.17 in /opt/anaconda3/lib/python3.9/site-packages (from arch) (1.20.3)
 Requirement already satisfied: statsmodels>=0.11 in /opt/anaconda3/lib/python3.9/site-packages (from arch) (0.12.2)
 Requirement already satisfied: pandas>=1.0 in /opt/anaconda3/lib/python3.9/site-packages (from arch) (1.3.4)
 Requirement already satisfied: scipy>=1.3 in /opt/anaconda3/lib/python3.9/site-packages (from arch) (1.7.1)
 Requirement already satisfied: python-dateutil>=2.7.3 in /opt/anaconda3/lib/python3.9/site-packages (from pandas>=1.0->arch) (2.8.2)
 Requirement already satisfied: pytz>=2017.3 in /opt/anaconda3/lib/python3.9/site-packages (from pandas>=1.0->arch) (2021.3)
 Requirement already satisfied: six>=1.5 in /opt/anaconda3/lib/python3.9/site-packages (from python-dateutil>=2.7.3->pandas>=1.0->arch) (1.16.0)
 Requirement already satisfied: patsy>=0.5 in /opt/anaconda3/lib/python3.9/site-packages (from statsmodels>=0.11->arch) (0.5.2)
 Installing collected packages: property-cached, arch
 Successfully installed arch-5.0.1 property-cached-1.6.4

In [162]: `stock_returns = prices[['dAIB', 'dBoI']].copy()`

In [163]: *# Detect ARCH effects in the data - AIB*

```
data_arch = sm.add_constant(stock_returns['dAIB'])
res_arch = sm.OLS(data_arch['dAIB'], data_arch['const']).fit()
print(res_arch.summary())
```

OLS Regression Results

```
=====
=====
Dep. Variable:          dAIB    R-squared:
-0.000
Model:                OLS    Adj. R-squared:
-0.000
Method:             Least Squares    F-statistic:
nan
Date:              Fri, 19 Nov 2021    Prob (F-statistic):
nan
Time:              10:46:57    Log-Likelihood:
-2117.9
No. Observations:      788    AIC:
4238.
```

```

Df Residuals:          787    BIC:
4242.
Df Model:              0
Covariance Type:      nonrobust

=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
const         -0.0936      0.127      -0.739      0.460      -0.342
0.155
=====
=====
Omnibus:          153.926    Durbin-Watson:
1.788
Prob(Omnibus):      0.000    Jarque-Bera (JB):
2722.788
Skew:              0.313    Prob(JB):
0.00
Kurtosis:          12.085    Cond. No.
1.00
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the error s is correctly specified.

/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/tsatools.py:142: FutureWarning:

In a future version of pandas all arguments of concat except for the argument 'objs' will be keyword-only

```

In [164]: resAIB = het_arch(res_arch.resid,5)
name = ['lm', 'lmpval', 'fval', 'fpval']
lzip(name,resAIB)

# p-value of the Le Grange Multiplier shows the AIB returns definit

```

```

Out[164]: [('lm', 179.37373729805307),
('lmpval', 7.280948968426872e-37),
('fval', 46.17870443765159),
('fpval', 8.130684806595673e-42)]

```

```

In [165]: # Detect ARCH effects in the data - BoI
data_arch2 = sm.add_constant(stock_returns['dBoI'])
res_arch2 = sm.OLS(data_arch2['dBoI'], data_arch2['const']).fit()
print(res_arch2.summary())

```

OLS Regression Results

```

=====
=====
Dep. Variable:                    dBoI    R-squared:
0.000

Model:                            OLS    Adj. R-squared:
0.000
Method:                        Least Squares    F-statistic:
nan
Date:                        Fri, 19 Nov 2021    Prob (F-statistic):
nan
Time:                        10:46:57    Log-Likelihood:
-2056.6
No. Observations:                788    AIC:
4115.
Df Residuals:                    787    BIC:
4120.
Df Model:                        0
Covariance Type:                nonrobust
=====
=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
const          -0.0454      0.117      -0.387      0.699      -0.276
0.185
=====
=====
Omnibus:                    82.059    Durbin-Watson:
1.937
Prob(Omnibus):                0.000    Jarque-Bera (JB):
549.364
Skew:                        0.124    Prob(JB):
5.09e-120
Kurtosis:                    7.083    Cond. No.
1.00
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the error
s is correctly specified.
/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/tsatool
s.py:142: FutureWarning:

In a future version of pandas all arguments of concat except for t
he argument 'objs' will be keyword-only

```
In [166]: resBoI = het_arch(res_arch2.resid,5)
name = ['lm', 'lmpval', 'fval', 'fpval']
lzip(name,resBoI)

# p-value of the Le Grange Multiplier shows the BoI returns definit
```

```
Out[166]: [('lm', 85.20678278339041),
 ('lmpval', 6.812577134688816e-17),
 ('fval', 18.97572764802119),
 ('fpval', 7.809618398613308e-18)]
```

```
In [167]: # GARCH model - AIB
from arch import arch_model

am = arch_model(stock_returns['dAIB'], vol='GARCH')
res = am.fit()
print(res.summary())

# Alpha & beta are both statisitcally significant
# We have come up with a model of volatility that is sufficient
```

```
Iteration:      1,   Func. Count:      6,   Neg. LLF: 4144.5300386
72261
Iteration:      2,   Func. Count:     14,   Neg. LLF: 5641441.8661
56031
Iteration:      3,   Func. Count:     20,   Neg. LLF: 3393.0743669
32486
Iteration:      4,   Func. Count:     26,   Neg. LLF: 1914.2272121
198616
Iteration:      5,   Func. Count:     32,   Neg. LLF: 1903.3433437
556175
Iteration:      6,   Func. Count:     37,   Neg. LLF: 1902.7538937
305376
Iteration:      7,   Func. Count:     42,   Neg. LLF: 2822.2556007
00125
Iteration:      8,   Func. Count:     49,   Neg. LLF: 1904.9451084
146092
Iteration:      9,   Func. Count:     55,   Neg. LLF: 1903.5571551
31903
Iteration:     10,   Func. Count:     61,   Neg. LLF: 1902.5703476
964213
Iteration:     11,   Func. Count:     66,   Neg. LLF: 1902.5702694
25189
Iteration:     12,   Func. Count:     71,   Neg. LLF: 1902.5701781
388946
Iteration:     13,   Func. Count:     76,   Neg. LLF: 1902.5701765
835415
Iteration:     14,   Func. Count:     80,   Neg. LLF: 1902.5701765
84205
Optimization terminated successfully      (Exit mode 0)
      Current function value: 1902.5701765835415
      Iterations: 14
      Function evaluations: 80
```

Gradient evaluations: 14

Constant Mean – GARCH Model Results

```

=====
=====
Dep. Variable:          dAIB    R-squared:
0.000
Mean Model:            Constant Mean    Adj. R-squared:
0.000
Vol Model:             GARCH    Log-Likelihood:
-1902.57
Distribution:          Normal    AIC:
3813.14
Method:               Maximum Likelihood    BIC:
3831.82
                                No. Observations:
788
Date:                 Fri, Nov 19 2021    Df Residuals:
787
Time:                 10:47:01    Df Model:
1

```

Mean Model

```

=====
=====

```

	coef	std err	t	P> t	95.0% Co
nf. Int.					
mu	-0.0965	7.580e-02	-1.273	0.203	[-0.245, 5.209e-02]

Volatility Model

```

=====
=====

```

	coef	std err	t	P> t	95.0% C
onf. Int.					
omega	0.1174	8.458e-02	1.388	0.165	[-4.839e-02, 0.283]
alpha[1]	0.1283	5.921e-02	2.167	3.025e-02	[1.224e-02, 0.244]
beta[1]	0.8694	5.430e-02	16.012	1.054e-57	[0.763, 0.976]

```

=====
=====

```

Covariance estimator: robust

In [168]: # E-GARCH – AIB

```

am_eg = arch_model(stock_returns['dAIB'], vol='EGARCH', o=1)
res_eg = am_eg.fit()
print(res_eg.summary())

# alpha gamma & beta all statistically significant

```

```

# alpha, gamma & beta are statistically significant
# The gamma term shows that not only for the volatility of AIB return
# but also some asymmetric affect between positive and negative vol.

Iteration:      1,   Func. Count:      7,   Neg. LLF: 1886586894.1
210165
Iteration:      2,   Func. Count:     17,   Neg. LLF: 4249.0868425
70632
Iteration:      3,   Func. Count:     27,   Neg. LLF: 127087373581
.51004
Iteration:      4,   Func. Count:     37,   Neg. LLF: 12888385617.
320312
Iteration:      5,   Func. Count:     46,   Neg. LLF: 1903.6713332
80587
Iteration:      6,   Func. Count:     53,   Neg. LLF: 425984074.75
55089
Iteration:      7,   Func. Count:     61,   Neg. LLF: 1903.2754498
366771
Iteration:      8,   Func. Count:     68,   Neg. LLF: 1899.6810227
391659
Iteration:      9,   Func. Count:     74,   Neg. LLF: 1899.6686394
735564
Iteration:     10,   Func. Count:     80,   Neg. LLF: 1899.6681482
66571
Iteration:     11,   Func. Count:     86,   Neg. LLF: 1899.6681429
520604
Iteration:     12,   Func. Count:     92,   Neg. LLF: 1899.6681416
80943
Iteration:     13,   Func. Count:     97,   Neg. LLF: 1899.6681416
803965
Optimization terminated successfully      (Exit mode 0)
      Current function value: 1899.668141680943
      Iterations: 13
      Function evaluations: 97
      Gradient evaluations: 13
      Constant Mean - EGARCH Model Results
=====
=====
Dep. Variable:                dAIB   R-squared:
0.000
Mean Model:                  Constant Mean   Adj. R-squared:
0.000
Vol Model:                   EGARCH   Log-Likelihood:
-1899.67
Distribution:                Normal   AIC:
3809.34
Method:                      Maximum Likelihood   BIC:
3832.68

                                     No. Observations:
788
Date:                        Fri, Nov 19 2021   Df Residuals:
787
Time:                        10:47:01   Df Model:
1

```


Mean Model					
	coef	std err	t	P> t	95.0% C
Conf. Int.					
mu	-0.1697	7.845e-02	-2.163	3.054e-02	[-0.323, -1.593e-02]
Volatility Model					
	coef	std err	t	P> t	95.0% C
Conf. Int.					
omega	0.0321	2.317e-02	1.385	0.166	[-1.333e-02, 7.751e-02]
alpha[1]	0.1915	9.673e-02	1.980	4.773e-02	[1.913e-03, 0.381]
gamma[1]	-0.0525	2.492e-02	-2.108	3.500e-02	[-0.101, -3.699e-03]
beta[1]	0.9892	8.971e-03	110.271	0.000	[0.972, 1.007]

Covariance estimator: robust

In [169]: *# The next graph was printing multiple pages of warnings so the fol*
Not recommended in practice, but for this purpose it should be ok

```
import warnings
warnings.filterwarnings("ignore")
```

In [170]: *# Finish by forecasting using GARCH – AIB*

Create two data samples

```
data_in_the_sample = stock_returns.loc[:'2019-08-05', 'dAIB']
data_out_of_the_sample = stock_returns.loc['2019-08-06':, 'dAIB']
```

Static forecasting

```
am = arch_model(stock_returns['dAIB'], vol='Garch')
cvar_dAIB_stat = {}
for date in data_out_of_the_sample.index:
    res = am.fit(last_obs=date, disp='off')
    forecasts = res.forecast(horizon=1)
    forecasts_res = forecasts.variance.dropna()
    cvar_dAIB_stat[date] = forecasts_res.iloc[1]
cvar_dAIB_stat = pd.DataFrame(cvar_dAIB_stat).T
```


Dynamic Forecasting

```
res = am.fit(last_obs= '2019-08-06', disp='off')
forecasts = res.forecast(horizon=len(data_out_of_the_sample))
forecasts_res = forecasts.variance.dropna()
cvar_dAIB_dyn = pd.DataFrame(data= forecasts_res.iloc[1].values,\
                             columns=['dynamic forecasting'],\
                             index=data_out_of_the_sample.index)
```

Now chart

```
plt.figure(1)
plt.plot(cvar_dAIB_stat, label = 'static forecast')
plt.plot(cvar_dAIB_dyn, label = 'dynamic forecast')
plt.xlabel('Date')
plt.ylabel('Forecasts')
plt.legend()
plt.show()

warnings.warn("first example of warning!", DeprecationWarning)
```



In [171]: *# Static forecast above is much more volatile*
It is not reasonable that we can predict the massive volatility
Dynamic is probably better for this time series

In [172]: *# GARCH model - BoI*
from arch **import** arch_model

am2 = arch_model(stock_returns['dBoI'], vol='GARCH')
res2 = am2.fit()
print(res2.summary())

Alpha & beta are both statisitcally significant
We have come up with a model that is sufficient

Iteration: 1, Func. Count: 6, Neg. LLF: 4144.5300386

```

72261
Iteration:      2,   Func. Count:      14,   Neg. LLF: 5641441.8661
56031
Iteration:      3,   Func. Count:      20,   Neg. LLF: 3393.0743669
32486
Iteration:      4,   Func. Count:      26,   Neg. LLF: 1914.2272121
198616
Iteration:      5,   Func. Count:      32,   Neg. LLF: 1903.3433437
556175
Iteration:      6,   Func. Count:      37,   Neg. LLF: 1902.7538937
305376
Iteration:      7,   Func. Count:      42,   Neg. LLF: 2822.2556007
00125
Iteration:      8,   Func. Count:      49,   Neg. LLF: 1904.9451084
146092
Iteration:      9,   Func. Count:      55,   Neg. LLF: 1903.5571551
31903
Iteration:     10,   Func. Count:      61,   Neg. LLF: 1902.5703476
964213
Iteration:     11,   Func. Count:      66,   Neg. LLF: 1902.5702694
25189
Iteration:     12,   Func. Count:      71,   Neg. LLF: 1902.5701781
388946
Iteration:     13,   Func. Count:      76,   Neg. LLF: 1902.5701765
835415
Iteration:     14,   Func. Count:      80,   Neg. LLF: 1902.5701765
84205

```

```

Optimization terminated successfully   (Exit mode 0)
      Current function value: 1902.5701765835415
      Iterations: 14
      Function evaluations: 80
      Gradient evaluations: 14

```

Constant Mean – GARCH Model Results

```

=====
=====
Dep. Variable:                  dAIB   R-squared:
0.000
Mean Model:                    Constant Mean   Adj. R-squared:
0.000
Vol Model:                     GARCH   Log-Likelihood:
-1902.57
Distribution:                   Normal   AIC:
3813.14
Method:                        Maximum Likelihood   BIC:
3831.82
                                No. Observations:
788
Date:                          Fri, Nov 19 2021   Df Residuals:
787
Time:                           10:47:15   Df Model:
1
                                Mean Model
=====
=====

```

	coef	std err	t	P> t	95.0% Co
nf. Int.					

mu	-0.0965	7.580e-02	-1.273	0.203	[-0.245, 5.209e-02]
Volatility Model					
=====					
=====					
onf. Int.	coef	std err	t	P> t	95.0% C

omega	0.1174	8.458e-02	1.388	0.165	[-4.839e-02, 0.283]
alpha[1]	0.1283	5.921e-02	2.167	3.025e-02	[1.224e-02, 0.244]
beta[1]	0.8694	5.430e-02	16.012	1.054e-57	[0.763, 0.976]
=====					
=====					

Covariance estimator: robust

In [173]: *# E-GARCH - BoI*

```
am_eg2 = arch_model(stock_returns['dBoI'], vol='EGARCH', o=1)
res_eg2 = am_eg2.fit()
print(res_eg2.summary())

# gamma & beta both statistically significant
# The gamma term shows that not only for the volatility of AIB return
# but also some asymmetric affect between positive and negative vol.
```

```
Iteration:      1,   Func. Count:      7,   Neg. LLF: 6929.1510965
33769
Iteration:      2,   Func. Count:     18,   Neg. LLF: 6236624667.3
22483
Iteration:      3,   Func. Count:     28,   Neg. LLF: 312014643647
.889
Iteration:      4,   Func. Count:     38,   Neg. LLF: 3284949847.6
14312
Iteration:      5,   Func. Count:     46,   Neg. LLF: 1918.6792505
186056
Iteration:      6,   Func. Count:     52,   Neg. LLF: 1974.5060670
68275
Iteration:      7,   Func. Count:     61,   Neg. LLF: 8914.0224633
8648
Iteration:      8,   Func. Count:     72,   Neg. LLF: 1919.3403510
141159
Iteration:      9,   Func. Count:     79,   Neg. LLF: 437347090.48
57292
Iteration:     10,   Func. Count:     86,   Neg. LLF: 2783616348.8
```

1074				
Iteration:	11,	Func. Count:	95,	Neg. LLF: 51612633077.
75505				
Iteration:	12,	Func. Count:	105,	Neg. LLF: 177886530632
9.2637				
Iteration:	13,	Func. Count:	115,	Neg. LLF: 474644985.89
41885				
Iteration:	14,	Func. Count:	123,	Neg. LLF: 131999712952
6.8967				
Iteration:	15,	Func. Count:	133,	Neg. LLF: 1912.9831260
28099				
Iteration:	16,	Func. Count:	140,	Neg. LLF: 3547.9718822
67667				
Iteration:	17,	Func. Count:	148,	Neg. LLF: 1906.0359251
761352				
Iteration:	18,	Func. Count:	154,	Neg. LLF: 142346695804
4.9404				
Iteration:	19,	Func. Count:	165,	Neg. LLF: 517392150974
7.839				
Iteration:	20,	Func. Count:	176,	Neg. LLF: 1095157.6992
387641				
Iteration:	21,	Func. Count:	183,	Neg. LLF: 1831627.9698
458614				
Iteration:	22,	Func. Count:	190,	Neg. LLF: 770127.01794
12183				
Iteration:	23,	Func. Count:	197,	Neg. LLF: 775628.61253
27232				
Iteration:	24,	Func. Count:	205,	Neg. LLF: 780050.46935
11531				
Iteration:	25,	Func. Count:	213,	Neg. LLF: 2935012.4677
572455				
Iteration:	26,	Func. Count:	221,	Neg. LLF: 1061935.6016
40977				
Iteration:	27,	Func. Count:	237,	Neg. LLF: 2856924.9780
078926				
Iteration:	28,	Func. Count:	250,	Neg. LLF: 2522636.0644
91747				
Iteration:	29,	Func. Count:	263,	Neg. LLF: 2510985.0695
961746				
Iteration:	30,	Func. Count:	279,	Neg. LLF: 2379791.5131
772477				
Iteration:	31,	Func. Count:	286,	Neg. LLF: 2425206.2516
00474				
Iteration:	32,	Func. Count:	294,	Neg. LLF: 1636262.8747
66716				
Iteration:	33,	Func. Count:	310,	Neg. LLF: 1597972.5947
601595				
Iteration:	34,	Func. Count:	326,	Neg. LLF: 1597252.0458
914507				
Iteration:	35,	Func. Count:	342,	Neg. LLF: 1597110.6325
41253				
Iteration:	36,	Func. Count:	358,	Neg. LLF: 1597023.0108
45362				
Iteration:	37,	Func. Count:	374,	Neg. LLF: 1601776.9565

507097

Iteration: 38, Func. Count: 390, Neg. LLF: 518761.64510
192344

Iteration: 39, Func. Count: 406, Neg. LLF: 145797257287
77.816

Optimization terminated successfully (Exit mode 0)
Current function value: 1902.2204114291446
Iterations: 40
Function evaluations: 416
Gradient evaluations: 39

Constant Mean – EGARCH Model Results

=====

Dep. Variable: dBoI R-squared:
0.000
Mean Model: Constant Mean Adj. R-squared:
0.000
Vol Model: EGARCH Log-Likelihood:
-1902.22
Distribution: Normal AIC:
3814.44
Method: Maximum Likelihood BIC:
3837.79
No. Observations:

788
Date: Fri, Nov 19 2021 Df Residuals:
787
Time: 10:47:16 Df Model:
1

Mean Model

=====

	coef	std err	t	P> t	95.0% Conf
. Int.					
mu	-0.2019	1.447e-04	-1395.082	0.000	[-0.202, -0.202]

Volatility Model

=====

	coef	std err	t	P> t	95.0
% Conf. Int.					
omega	6.7477e-03	8.617e-10	7.831e+06	0.000	[6.748e-03, 6.748e-03]
alpha[1]	-0.0283	6.563e-05	-430.887	0.000	[-2.841e-02, -2.815e-02]
gamma[1]	-0.0619	6.302e-05	-982.243	0.000	[-6.202e-02, -6.178e-02]
beta[1]	0.9980	1.812e-09	5.507e+08	0.000	[0.998, 0.998]

=====

Covariance estimator: robust

```

In [174]: # Finish by forecasting using GARCH – BoI

# Create two data samples

data_in_the_sample = stock_returns.loc[:'2019-08-05', 'dBoI']
data_out_of_the_sample = stock_returns.loc['2019-08-06':, 'dBoI']

# Static forecasting

am = arch_model(stock_returns['dBoI'], vol='Garch')
cvar_dBoI_stat = {}
for date in data_out_of_the_sample.index:
    res = am.fit(last_obs=date, disp='off')
    forecasts = res.forecast(horizon=1)
    forecasts_res = forecasts.variance.dropna()
    cvar_dBoI_stat[date] = forecasts_res.iloc[1]
cvar_dBoI_stat = pd.DataFrame(cvar_dBoI_stat).T


# Dynamic Forecasting

res = am.fit(last_obs= '2019-08-06', disp='off')
forecasts = res.forecast(horizon=len(data_out_of_the_sample))
forecasts_res = forecasts.variance.dropna()
cvar_dBoI_dyn = pd.DataFrame(data= forecasts_res.iloc[1].values,\
                             columns=['dynamic forecasting'],\
                             index=data_out_of_the_sample.index)

# Now chart

plt.figure(1)
plt.plot(cvar_dBoI_stat, label = 'static forecast')
plt.plot(cvar_dBoI_dyn, label = 'dynamic forecast')
plt.xlabel('Date')
plt.ylabel('Forecasts')
plt.legend()
plt.show()

```




```
In [175]: # Static forecast above is much more volatile
# Not reasonable that we can predict the massive volatility
# Dynamic is probably better for this time series
```

Part 2(c) VAR Model

2C: Develop a VAR model of the two series and show how one series might influence the other

```
In [176]: prices.corr()

# 98.4% correlation between the stock prices
# 73.6% correlation between the stock returns
```

Out[176]:

	AIB	Bol	dAIB	dBol
AIB	1.000000	0.983994	0.004356	-0.007269
Bol	0.983994	1.000000	0.002867	0.006741
dAIB	0.004356	0.002867	1.000000	0.736466
dBol	-0.007269	0.006741	0.736466	1.000000

```
In [177]: stock_returns = prices[['dAIB', 'dBoI']].copy()
```

```
In [178]: from statsmodels.tsa.vector_ar.vecm import coint_johansen
```

In [179]: *# Test for cointegration of the stock returns*

```
def cointegration_test(stock_returns, alpha=0.05):
    """Perform Johanson's Cointegration Test and Report Summary"""
    out = coint_johansen(stock_returns, -1, 5)
    d = {'0.90':0, '0.95':1, '0.99':2}
    traces = out.lr1
    cvts = out.cvt[:, d[str(1-alpha)]]
    def adjust(val, length= 6): return str(val).ljust(length)

    # Summary
    print('Name    :: Test Stat > C(95%)    => Signif \n', '--'*length)
    for col, trace, cvt in zip(stock_returns.columns, traces, cvts):
        print(adjust(col), ':: ', adjust(round(trace,2), 9), ">", adjust(cvt, 9))

cointegration_test(stock_returns)

# Cointegration is evident
```

Name	::	Test Stat > C(95%)	=>	Signif
dAIB	::	298.72 > 12.3212	=>	True
dBoI	::	117.12 > 4.1296	=>	True

In [180]: *# Split the dataset into train and test data*
VAR model will forecast next year's values

```
nobs = 260
stock_returns_train, stock_returns_test = stock_returns[1:-nobs], stock_returns[-nobs:]

# Check size
print(stock_returns_train.shape)
print(stock_returns_test.shape)

(527, 2)
(260, 2)
```

In [181]: *# Selecting Order(p) of the VAR model*
Look for order that gives the lowest AIC

```
from statsmodels.tsa.api import VAR

model = VAR(stock_returns)
for i in [1,2,3,4,5,6,7,8,9]:
    result = model.fit(i)
    print('Lag Order =', i)
    print('AIC : ', result.aic)
    print('BIC : ', result.bic)
    print('FPE : ', result.fpe)
    print('HQIC: ', result.hqic, '\n')

# Lowest AIC at lag 6
```

Lag Order = 1
AIC : 4.136117017390967
BIC : 4.1717070675695
FPE : 62.5594366503173
HQIC: 4.14979967347309

Lag Order = 2
AIC : 4.1399434192711375
BIC : 4.199319459887286
FPE : 62.799289675230334
HQIC: 4.162772007128729

Lag Order = 3
AIC : 4.14411857578782
BIC : 4.227328221710054
FPE : 63.06207263553421
HQIC: 4.176112500484777

Lag Order = 4
AIC : 4.129890920347703
BIC : 4.236981943774116
FPE : 62.17126637119178
HQIC: 4.171069652619446

Lag Order = 5
AIC : 4.123336373176197
BIC : 4.2543567043534605
FPE : 61.765198233386876
HQIC: 4.173719449733416

Lag Order = 6
AIC : 4.12180864922691
BIC : 4.276806377173429
FPE : 61.671059921409025
HQIC: 4.181415673061101

Lag Order = 7
AIC : 4.125071632974485
BIC : 4.304095006207763
FPE : 61.87282577708293
HQIC: 4.19392227366527

Lag Order = 8
AIC : 4.13121922085623
BIC : 4.334316648124491
FPE : 62.254637383354215
HQIC: 4.209333214880486

Lag Order = 9
AIC : 4.1308648676184125
BIC : 4.358084918636497
FPE : 62.23292623332048
HQIC: 4.218262018661224

In [182]: *# Training the VAR model of selected Order(p)*

```
model_fitted = model.fit(6)
model_fitted.summary()
```

Out[182]: Summary of Regression Results

```
=====
Model:                                VAR
Method:                               OLS
Date:                                Fri, 19, Nov, 2021
Time:                                10:47:45
-----
--
No. of Equations:                    2.00000    BIC:                                4.276
81
Nobs:                                782.000    HQIC:                               4.181
42
Log likelihood:                      -3804.85    FPE:                                61.67
11
AIC:                                4.12181    Det(Omega_mle):                     59.67
06
-----
--
Results for equation dAIB
=====
=====
```

	coefficient	std. error	t-stat
prob			

const	-0.120742	0.125740	-0.960
0.337			
L1.dAIB	0.067411	0.053538	1.259
0.208			
L1.dBoI	0.022445	0.057130	0.393
0.694			
L2.dAIB	-0.107096	0.053337	-2.008
0.045			
L2.dBoI	0.084193	0.057158	1.473
0.141			
L3.dAIB	-0.109968	0.053005	-2.075
0.038			
L3.dBoI	0.085928	0.057210	1.502
0.133			
L4.dAIB	-0.217026	0.053031	-4.092
0.000			
L4.dBoI	0.197156	0.057232	3.445
0.001			
L5.dAIB	-0.134342	0.053519	-2.510
0.012			
L5.dBoI	0.168451	0.057437	2.933

```

0.003
L6.dAIB      -0.076845      0.053476      -1.437
0.151
L6.dBoI      0.020261      0.057546      0.352
0.725

```

Results for equation dBoI

```

=====
=====

```

	coefficient	std. error	t-stat
prob			

const	-0.068848	0.117641	-0.585
0.558			
L1.dAIB	0.115824	0.050090	2.312
0.021			
L1.dBoI	-0.067220	0.053450	-1.258
0.209			
L2.dAIB	-0.022980	0.049901	-0.460
0.645			
L2.dBoI	-0.022610	0.053477	-0.423
0.672			
L3.dAIB	-0.054121	0.049591	-1.091
0.275			
L3.dBoI	0.005093	0.053525	0.095
0.924			
L4.dAIB	-0.078010	0.049615	-1.572
0.116			
L4.dBoI	0.102353	0.053546	1.911
0.056			
L5.dAIB	-0.056160	0.050072	-1.122
0.262			
L5.dBoI	0.039339	0.053738	0.732
0.464			
L6.dAIB	-0.112538	0.050032	-2.249
0.024			
L6.dBoI	0.113613	0.053840	2.110
0.035			

```

=====
=====

```

Correlation matrix of residuals

```

          dAIB      dBoI
dAIB    1.000000  0.741030
dBoI    0.741030  1.000000

```

Or Alternatively:

In [183]: `# Run VAR – for now assume we know the appropriate lag length is 5`

```

stock_returns = prices[['dAIB', 'dBoI']].copy()

model = smt.VAR(stock_returns)
res=model.fit(maxlags=5)
print(res.summary())

# For AIB: lag 4 and lag 5 are significant for both
# For BoI: Influenced by AIB's previous day returns

```

Summary of Regression Results

```

=====
Model:                                VAR
Method:                               OLS
Date:                                Fri, 19, Nov, 2021
Time:                                10:47:45
-----
--
No. of Equations:                     2.00000    BIC:                                4.254
36
Nobs:                                783.000    HQIC:                               4.173
72
Log likelihood:                       -3814.34    FPE:                                61.76
52
AIC:                                  4.12334    Det(Omega_mle):                     60.06
57
-----
--

```

Results for equation dAIB

```

=====
=====

```

	coefficient	std. error	t-stat
prob			

const	-0.114026	0.125602	-0.908
0.364			
L1.dAIB	0.070008	0.053066	1.319
0.187			
L1.dBoI	0.022978	0.056907	0.404
0.686			
L2.dAIB	-0.094915	0.052756	-1.799
0.072			
L2.dBoI	0.075375	0.056935	1.324
0.186			
L3.dAIB	-0.105809	0.052856	-2.002
0.045			
L3.dBoI	0.084765	0.057045	1.486
0.137			
L4.dAIB	-0.212242	0.052782	-4.021
0.000			
L4.dBoI	0.194266	0.056962	3.410
0.001			
L5.dAIB	-0.137058	0.053187	-2.577

```

0.010
L5.dBoI      0.163066      0.057023      2.860
0.004
=====
=====

```

Results for equation dBoI

```

=====
=====

```

	coefficient	std. error	t-stat
prob			

const	-0.060463	0.117679	-0.514
0.607			
L1.dAIB	0.129913	0.049719	2.613
0.009			
L1.dBoI	-0.075804	0.053317	-1.422
0.155			
L2.dAIB	-0.006812	0.049428	-0.138
0.890			
L2.dBoI	-0.031608	0.053344	-0.593
0.553			
L3.dAIB	-0.046207	0.049522	-0.933
0.351			
L3.dBoI	-0.001408	0.053447	-0.026
0.979			
L4.dAIB	-0.067142	0.049453	-1.358
0.175			
L4.dBoI	0.092031	0.053369	1.724
0.085			
L5.dAIB	-0.049210	0.049832	-0.988
0.323			
L5.dBoI	0.025736	0.053426	0.482
0.630			

```

=====
=====

```

Correlation matrix of residuals

```

          dAIB      dBoI
dAIB      1.000000  0.740115
dBoI      0.740115  1.000000

```

In [184]: *# Test up to 10 lags to find out the appropriate lag length*

```
res=model.select_order(maxlags=10)
print(res.summary())
```

```
# For AIC the min is 10
# For BIC the min is 0
# For FPE the min is 10
# For HQIC the min is 0
```

VAR Order Selection (* highlights the minimums)

	AIC	BIC	FPE	HQIC
0	4.161	4.173*	64.13	4.166*
1	4.153	4.189	63.63	4.167
2	4.156	4.215	63.79	4.179
3	4.159	4.243	64.03	4.192
4	4.143	4.251	62.99	4.184
5	4.134	4.265	62.40	4.184
6	4.130	4.286	62.17	4.190
7	4.131	4.311	62.24	4.200
8	4.135	4.338	62.48	4.213
9	4.133	4.360	62.34	4.220
10	4.126*	4.377	61.91*	4.222

In [185]: *# Go with AIC (as HQIC is 0)*

```
model_2 = smt.VAR(stock_returns)
res_2 = model_2.fit(maxlags=10)
print(res_2.summary())
```

```
# Lags 4 and 5 are relevant for AIB
# Lags 4, 6 and 10 are relevant for BoI, the return from 4, 6 and 10
```

Summary of Regression Results

=====			
Model:	VAR		
Method:	OLS		
Date:	Fri, 19, Nov, 2021		
Time:	10:47:45		

--			
No. of Equations:	2.00000	BIC:	4.377
04			
Nobs:	778.000	HQIC:	4.222
35			
Log likelihood:	-3770.75	FPE:	61.90
96			
AIC:	4.12565	Det(Omega_mle):	58.69
81			

--

Results for equation dAIB

=====			
=====			
prob	coefficient	std. error	t-stat

const	-0.092336	0.124672	-0.741
0.459			
L1.dAIB	0.070150	0.053539	1.310
0.190			
L1.dBoI	0.018271	0.056756	0.322
0.748			
L2.dAIB	-0.119778	0.053523	-2.238
0.025			
L2.dBoI	0.090819	0.056904	1.596
0.110			
L3.dAIB	-0.113385	0.053706	-2.111
0.035			
L3.dBoI	0.082741	0.057134	1.448
0.148			
L4.dAIB	-0.229879	0.053663	-4.284
0.000			
L4.dBoI	0.222617	0.057312	3.884
0.000			
L5.dAIB	-0.122914	0.054319	-2.263
0.024			
L5.dBoI	0.166667	0.057637	2.892
0.004			
L6.dAIB	-0.069539	0.054107	-1.285
0.199			
L6.dBoI	0.018931	0.057893	0.327
0.744			
L7.dAIB	0.022045	0.053721	0.410
0.682			
L7.dBoI	0.100532	0.058003	1.733
0.083			
L8.dAIB	-0.096691	0.053523	-1.807
0.071			
L8.dBoI	0.046617	0.057931	0.805
0.421			
L9.dAIB	0.093287	0.053429	1.746
0.081			
L9.dBoI	-0.001737	0.057707	-0.030
0.976			
L10.dAIB	0.168958	0.053236	3.174
0.002			
L10.dBoI	-0.067641	0.057364	-1.179
0.238			
=====			
=====			

Results for equation dBoI

prob	coefficient	std. error	t-stat
const	-0.053831	0.117997	-0.456
0.648			
L1.dAIB	0.116249	0.050672	2.294
0.022			
L1.dBoI	-0.068516	0.053717	-1.275
0.202			
L2.dAIB	-0.032239	0.050657	-0.636
0.525			
L2.dBoI	-0.015288	0.053858	-0.284
0.777			
L3.dAIB	-0.052540	0.050831	-1.034
0.301			
L3.dBoI	-0.001579	0.054076	-0.029
0.977			
L4.dAIB	-0.085050	0.050790	-1.675
0.094			
L4.dBoI	0.118405	0.054243	2.183
0.029			
L5.dAIB	-0.049229	0.051411	-0.958
0.338			
L5.dBoI	0.037166	0.054551	0.681
0.496			
L6.dAIB	-0.104196	0.051210	-2.035
0.042			
L6.dBoI	0.108366	0.054794	1.978
0.048			
L7.dAIB	0.014815	0.050845	0.291
0.771			
L7.dBoI	0.062294	0.054897	1.135
0.256			
L8.dAIB	-0.042943	0.050658	-0.848
0.397			
L8.dBoI	-0.003253	0.054830	-0.059
0.953			
L9.dAIB	0.035284	0.050568	0.698
0.485			
L9.dBoI	0.014475	0.054618	0.265
0.791			
L10.dAIB	0.122223	0.050386	2.426
0.015			
L10.dBoI	-0.052077	0.054293	-0.959
0.337			

Correlation matrix of residuals

	dAIB	dBoI
dAIB	1.000000	0.736317

dBoI 0.736317 1.000000

In [186]: *# Granger Causality Test*

```

maxlag = 6

test = 'ssr_chi2test'
def grangers_causation_matrix(data, variables, test='ssr_chi2test',
    """Check Granger Causality of all possible combinations of the Time Series
    The rows are the response variable, columns are predictors. The values in the
    are the P-Values. P-Values lesser than the significance level (0.05) indicate
    the Null Hypothesis that the coefficients of the corresponding predictor variables
    zero, that is, the x does not cause Y can be rejected.

    data      : pandas dataframe containing the time series variables
    variables  : list containing names of the time series variables.
    """

    df = pd.DataFrame(np.zeros((len(variables), len(variables))), columns=variables)
    for c in df.columns:
        for r in df.index:
            test_result = grangercausalitytests(data[[r,c]], maxlag=maxlag)
            p_values = [round(test_result[i+1][0][test][1],4) for i in range(len(test_result)-1)]
            if verbose: print(f'Y = {r}, X = {c}, P Values = {p_values}')
            min_p_value = np.min(p_values)
            df.loc[r,c] = min_p_value
    df.columns = [var + '_x' for var in variables]
    df.index = [var + '_y' for var in variables]
    return df

gc = grangers_causation_matrix(stock_returns, variables = stock_returns.columns)
gc

# If p-value is < significance level of 0.05, then the corresponding variable is a significant
# Both AIB and BoI returns definitely have a significant effect on each other
# They cause effects on the other
# This causality makes these time series appropriate for a VAR model

```

Out[186]:

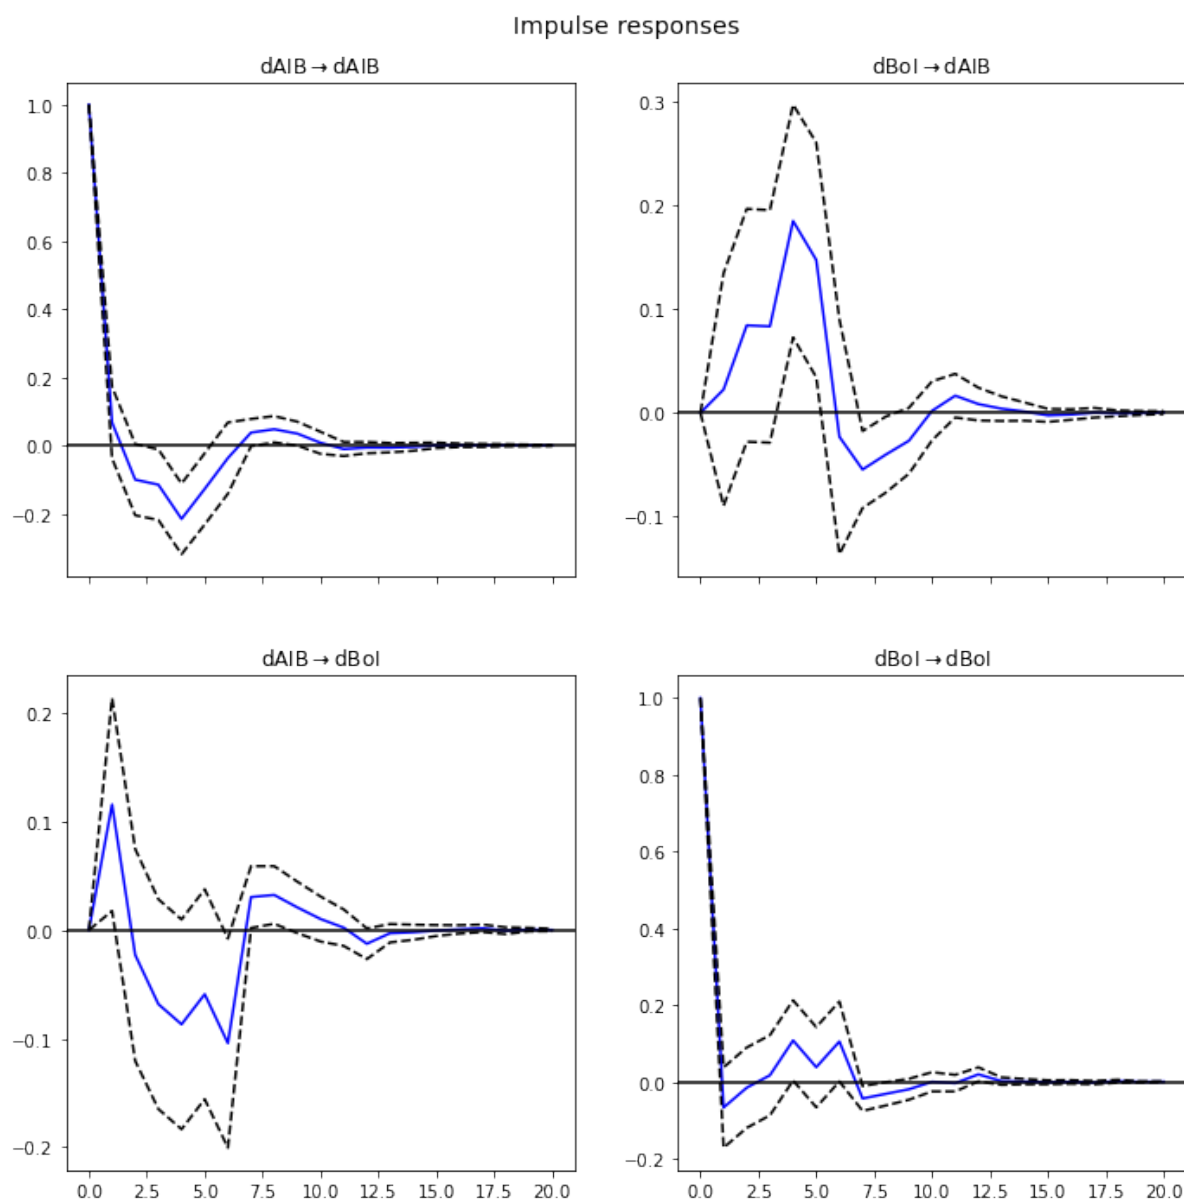
	dAIB_x	dBoI_x
dAIB_y	1.0000	0.0012
dBoI_y	0.0052	1.0000

In [187]: *# Impulse Response Functions*

```
model_ir = smt.VAR(stock_returns)
res_ir = model_ir.fit(maxlags=6)
```

```
irf=res_ir.irf(20)
irf.plot();
```

A negative shock in AIB, is likely to cause a negative return to I
And vice-versa



Part 3 - Machine Learning House Prices

Q3: Use appropriate machine learning techniques to develop a model to explain house prices

```
In [188]: houses = pd.read_csv('A1HousePrices.csv')
```

```
In [189]: houses.head(3)
```

```
Out[189]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floo
0	7129300520	20141013T000000	221900.0	3	1.00	1180	5650	1
1	6414100192	20141209T000000	538000.0	3	2.25	2570	7242	2
2	5631500400	20150225T000000	180000.0	2	1.00	770	10000	1

3 rows x 21 columns

```
In [190]: houses.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    21613 non-null  int64
1   date                  21613 non-null  object
2   price                 21613 non-null  float64
3   bedrooms              21613 non-null  int64
4   bathrooms             21613 non-null  float64
5   sqft_living           21613 non-null  int64
6   sqft_lot              21613 non-null  int64
7   floors                21613 non-null  float64
8   waterfront            21613 non-null  int64
9   view                  21613 non-null  int64
10  condition              21613 non-null  int64
11  grade                 21613 non-null  int64
12  sqft_above            21613 non-null  int64
13  sqft_basement         21613 non-null  int64
14  yr_built              21613 non-null  int64
15  yr_renovated          21613 non-null  int64
16  zipcode               21613 non-null  int64
17  lat                   21613 non-null  float64
18  long                  21613 non-null  float64
19  sqft_living15         21613 non-null  int64
20  sqft_lot15            21613 non-null  int64
dtypes: float64(5), int64(15), object(1)
memory usage: 3.5+ MB
```

```
In [191]: conv_dates = [1 if values == 2014 else 0 for values in houses.date]
houses['date'] = conv_dates
houses = houses.drop(["id", "sqft_living15", "sqft_lot15", "sqft_above15"])
houses.head(3)

# Dropping variables that are too unique, have too many missing values
# Sqft_living is a sum of sqft_above and sqft_basement. And basement
```

Out[191]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition
0	0	221900.0	3	1.00	1180	5650	1.0	0	0	0
1	0	538000.0	3	2.25	2570	7242	2.0	0	0	0
2	0	180000.0	2	1.00	770	10000	1.0	0	0	0

```
In [192]: houses.info()
```

```
# No missing values
# No strings
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 15 columns):
#   Column          Non-Null Count  Dtype
---  -
0   date            21613 non-null  int64
1   price           21613 non-null  float64
2   bedrooms        21613 non-null  int64
3   bathrooms        21613 non-null  float64
4   sqft_living      21613 non-null  int64
5   sqft_lot         21613 non-null  int64
6   floors           21613 non-null  float64
7   waterfront       21613 non-null  int64
8   view            21613 non-null  int64
9   condition        21613 non-null  int64
10  grade            21613 non-null  int64
11  yr_built         21613 non-null  int64
12  zipcode          21613 non-null  int64
13  lat              21613 non-null  float64
14  long             21613 non-null  float64
dtypes: float64(5), int64(10)
memory usage: 2.5 MB
```

```
In [193]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

```
In [194]: reg = LinearRegression()
```

```
In [195]: # Set the labels as the price column as prices are to be predicted

labels = houses['price']
train1 = houses.drop(['price'],axis=1)
```

```
In [196]: from sklearn.model_selection import train_test_split
```

```
In [197]: # 80% train data, 20% test data

x_train , x_test , y_train , y_test = train_test_split(train1 , labels ,
                                                    test_size=0.2,
                                                    random_state=42)
x_train.shape, y_train.shape, x_test.shape, y_test.shape
```

```
Out[197]: ((17290, 14), (17290,), (4323, 14), (4323,))
```

```
In [198]: reg.fit(x_train, y_train)
```

```
Out[198]: LinearRegression()
```

```
In [199]: # After fitting the model the score is 71.3%

reg.score(x_test, y_test)
```

```
Out[199]: 0.7131150403769968
```

Gradient Boosting

We can try to improve this 71.3% with gradient boosting

```
In [200]: from sklearn import ensemble
clf = ensemble.GradientBoostingRegressor(n_estimators = 400, max_depth=5,
                                         learning_rate = 0.1, loss = 'ls')
```

```
In [201]: clf.fit(x_train, y_train)
```

```
Out[201]: GradientBoostingRegressor(max_depth=5, n_estimators=400)
```

```
In [202]: clf.score(x_test,y_test)

# Accuracy has improved to 89.2%, a big improvement
# For weak predictions, gradient boosting works really well, even u
```

```
Out[202]: 0.8916183345372595
```

Alternatively, a decision tree could have been used

In [203]: *# Decision Tree*

```
dt = DecisionTreeClassifier()

dt.fit(x_train, y_train)
y_pred = dt.predict(x_test)
print("Accuracy is ", accuracy_score(y_test, y_pred)*100)
```

Accuracy is 0.9715475364330326

In [204]: *# Visualise the decision tree*

```
plt.figure(figsize=(16,16))
plot_tree(dt, max_depth=3, fontsize=10, feature_names=x_train.columns)
plt.show()
```

*# First step: Is the latitude of the house less than or equal to 47
 # Then is the sqft_living less than or equal to 1935?
 # And so on
 # So, location and then living area were the first two questions*

